

OurSys: An Interactive Debugger for SQL

ABSTRACT

SQL is declarative in nature and rich in its features. Writing semantically correct SQL queries and finding logical bugs in SQL are not easy, even for experienced programmers, who are often used to the mindset of working with a variety of general-purpose programming languages (GPLs). While there are many tools to assist with GPL debugging, SQL debugging has received much less attention. In this paper, we present OurSys, a SQL debugger that enables users to inspect the logical execution of SQL queries visually and interactively to identify and potentially fix logical bugs in the queries. OurSys draws analogies to the debugging paradigm of GPLs (e.g., stepping, watchpoints, etc.), making it easier for programmers to adopt. However, unlike debugging GPLs, which involves executing the underlying program in full to the point of interest, OurSys allows users to jump to arbitrary points of interest by leveraging the power of the database systems, through selective materialization and query rewrites. To simplify deployment, OurSys acts as a lightweight middleware on top of the database system; it imposes no overhead to prepare a database for debugging and maintains no state in the database systems during debugging sessions. We demonstrate the effectiveness of OurSys through performance experiments as well as a user study in an educational setting.

1 INTRODUCTION

Relational databases form the backbone of many data-intensive applications and scalable data analytics. Despite its age, SQL continues to retain its prevalence and importance due to its highly *declarative* nature (i.e., specifying what the answer should be rather than how to compute it) and its extensive set of features that have grown over time. On the other hand, SQL is difficult to understand and debug. In debugging GPLs (such as C++ or Python) that are typically *procedural* (i.e. explicitly describing the processing steps required to compute the answer), it is natural to trace the execution of programs to debug them. However, this method becomes much trickier for SQL.

As a first attempt, one may consider “tracing” a query’s logical or physical plan, a tree whose leaves represent base tables and internal nodes represent relational operators. Through this approach, a user can examine the intermediate results produced by each of the plan nodes. Unfortunately, there are several problems with this approach. Firstly, the database optimizer often compiles a SQL query into a plan that bears no resemblance to the original query, making plan tracing unhelpful in finding and fixing logical bugs in the original query. Secondly, debugging is usually iterative: the user may examine execution multiple times, sometimes with minor modifications to the query. However, even small changes can lead to a very different plan (e.g., addition or removal of a condition in WHERE can enable or disable an index scan opportunity). Even if the physical plan remains the same, there is no guarantee that the execution and result order are reproducible. For example, size of buffer memory, choice of hash function, and variations in the speed

of parallel threads at runtime can change the ordering of intermediate result rows. Such non-repeatable and seemingly inconsistent behaviors significantly complicate debugging.

Perhaps, one possible workaround would be to restrict the database optimizer to avoid optimization across syntactic blocks of a query, such that each subquery corresponds to some subtree in the plan, and the user can at least inspect the result of each subquery. However, correlated subqueries, a SQL construct frequently used in practice, render this workaround ineffectual for many queries. In Section 2, we give concrete examples illustrating this challenge and showing how OurSys aids in debugging these types of queries.

In this work, we focus on finding *logical errors* instead of fixing *performance issues*, and we aim to build an interactive SQL debugger with the following desiderata:

- (1) The debugger should conceptually execute a SQL query in a manner that is completely reproducible, faithful to how it is written, and easy for average programmers to understand.
- (2) The debugger should offer features analogous to those already in GPL debuggers, so they are easy to learn and adopt.
- (3) Unlike GPL debuggers, which usually must execute underlying programs in full in order to reach points of interest, this debugger should leverage the power of SQL and database systems to support more powerful features and/or more efficient implementation of features.
- (4) The debugger must scale to large database instances and gracefully handle prohibitively large intermediate results.
- (5) The debugger should be simple to run (e.g., from a remote browser) and easy to deploy on top of a database, without modifying database system internals or requiring extensive preparation of the database for debugging.

At a first glance, (1) necessitates executing the SQL query “literally” using a completely unoptimized plan, which runs counter to (4). Our key insight here, however, is that at any given point in time, the user can be examining only a small “window” of the entire execution. It suffices to support fast access to any given “window” without incurring the full cost of the unoptimized plan. Supporting such accesses, along with (2) and (3) while respecting (5), require novel optimization ideas — simply relying on SQL’s built-in OFFSET and LIMIT operators fails to deliver acceptable performance for interactive debugging.

Addressing the above challenges through the development of OurSys, we make the following contributions:

- OurSys introduces a novel debugging paradigm for SQL that draws many parallels to GPL debugging, where each query block is viewed as a function and correlated subqueries can be considered functions with arguments. We define the *canonical execution* of a SQL query, which is reproducible and faithful to query syntax; edits to a query lead to predictable changes in its execution.
- OurSys supports various debugging features analogous to GPL debugging, including stepping through execution, pausing to examine a particular point during execution (breakpoints), pausing automatically at points of interest (watchpoints), drilling down

into subqueries (stepping “into” a function), and row-level tracing (information flow analysis) in both forward (from input to output) and backward (from output to input) directions.

- While performing the canonical execution would have been extremely inefficient and impractical, OURSYS supports effectively fast “teleporting” from one point of execution to another without paying the cost of execution in between. We do so by constructing queries that directly compute, for a given window, information needed for debugging, and by applying selective materialization and rewrite optimizations to ensure their efficiency on large database instances.
- OURSYS has a web-based frontend and a middleware backend that runs on top of a database system. It imposes no overhead to prepare a database for debugging and maintains no state in the database during active debugging sessions. This architecture makes OURSYS easy to adopt and deploy.
- Our performance evaluation using the TPC-H benchmark shows the scalability of OURSYS on large databases. More specifically, it also demonstrates the advantage of our optimization techniques over standard database (PostgreSQL) support for retrieving windows of query result rows.
- We evaluate the efficacy of OURSYS with a large-scale user study of over 100 students in an introductory-level database course. Its findings indicate that OURSYS significantly improves students’ efficiency in debugging SQL queries.

2 EXAMPLE USE OF OURSYS FOR DEBUGGING

Since OURSYS conceptually executes SQL queries the way they are written, it is particularly suitable for novices who are learning how SQL queries work on the logical level. Furthermore, it serves as a powerful tool for novices and data professionals alike to find and fix logical bugs. This section provides a walk-through of how to use OURSYS to debug an incorrect query; we introduce the interface, concepts, and features of OURSYS. More formal and detailed discussions will be presented in later sections.

Example 2.1. Consider the toy database in Figure 1, which stores information about beers, bars serving them, and drinkers who like beers and frequent bars. We want to write a query for the following task: suppose every time a drinker frequents a bar, they buy one bottle of every beer they like or any beer priced \$2 or lower; find the expected weekly revenue of each bar and rank them by revenue from high to low. A user may come up with the following (incorrect) query:

```
|| SELECT s.bar, SUM(f.times_a_week * s.price) AS revenue -- Q
|| FROM Serves s, Frequents f
|| WHERE f.bar = s.bar
|| AND (s.price <= 2 OR
||      EXISTS (
||        SELECT * FROM Likes l WHERE f.drinker = l.drinker -- Qinner
||      ))
|| GROUP BY s.bar;
```

The above query intends to first find drinkers and beers available for purchase using a join between Serves and Frequents. It additionally applies the two (alternative) conditions for purchase: 1) the price is lower than \$2 and 2) the beer is liked by the drinker. Then, the query groups the intermediate results by bar and calculates the sum of revenue. There is a bug in the EXISTS subquery Q_{inner} , but the question for now is: how would a user examine the result of Q_{inner} ?

bar	beer	price
Apex	Corona	1
Apex	Dixie	2
Edge	Amstel	4
Edge	Corona	1.5
Tavern	Amstel	3
Tavern	Erdinger	1

(a) Serves

drinker	bar	times
Amy	Apex	1
Ben	Edge	4
Coy	Tavern	2
Dan	Edge	3

(b) Frequents

drinker	beer
Amy	Erdinger
Ben	Budweiser
Ben	Dixie
Coy	Amstel
Dan	Amstel
Dan	Corona

(c) Likes

Figure 1: A toy database about beers, bars, and drinkers.

Note that Q_{inner} is correlated, with the value for $f.drinker$ coming from the outer (i.e. enclosing) block. As a result, there is no way to inspect this result independently. This situation cannot be handled by a query plan with relational operators, where the result of each subtree depends on this subtree alone.

Indeed, most database optimizers will rewrite the above query for execution such that the subquery is decorrelated. The decorrelated subquery would be a join involving both Likes and Frequents to compute, effectively, the original subquery for all possible drinker values in a single effort. Then, the result will be further combined with the rest of the outer query. The new plan now consists of only relational operators and can be computed/debugged in a bottom-up fashion, but unfortunately, it is vastly different to the original query. Any user without in-depth knowledge of query optimization will likely be very confused.

Example 2.2. This incorrect query returns the following result for the database in Figure 1:

bar	revenue
Apex	3
Edge	38.5
Tavern	8

Based on the user’s knowledge of the database instance, the revenue of bar Edge seems to be higher than expected. We now walk through how to use OURSYS to debug the query, starting with this observation.

OURSYS presents a panel of UI debugging elements for each block of the query. Figure 2 illustrates the UI for the outer query block Q (for simplicity, we do not show the actual interface here as it contains other details that may be distracting for this discussion). OURSYS shows the execution of this block in stages, from top to bottom. At the very top, OURSYS shows all input tables in FROM. Note that one row from each input table is highlighted; this input combination (“combo”) intuitively defines the current point of execution being examined. Then, OURSYS presents the “joined & filtered” result, which is the intermediate output after the WHERE clause is applied. The intermediate result row produced by the current input combo is automatically highlighted for users. Between this result table and the input tables, a “filter expression” tree shows how the WHERE condition evaluates over the current input combo. The user can examine the value of each subexpression therein and see how the truth values (color-coded here) are combined by logical connectives. Following the joined & filtered result, OURSYS shows the GROUP BY result. For each group, in addition to the GROUP BY value, each group member’s contribution to the final SUM aggregate is also shown. Again, the group and the member that the current input combo contributes to are automatically highlighted. Finally, the final result of the query block is shown as the output table.

For the convenience of subsequent discussion, we show a symbolic row identifier (e.g. $s_0, f_2, j_6...$) for each row in all tables. We do in fact assign internal row identifiers, whose purposes will be explained later in Section 3, but they are not explicitly displayed by the UI.

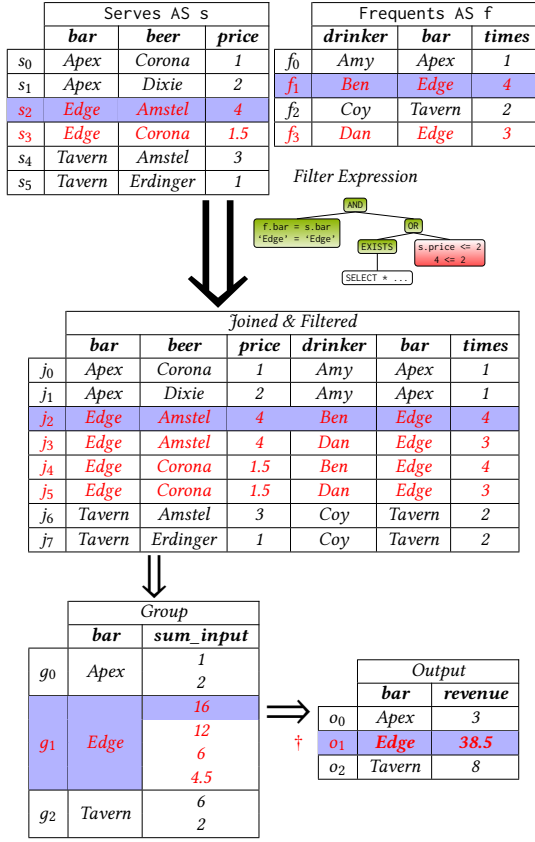


Figure 2: Debugging context for outer query block, Example 2.2.




Figure 3: Debugging context for inner query block, Example 2.1.

Given the unexpectedly high revenue of Edge, the user naturally wants to examine how that particular output row was computed. OURSYS supports tracing backwards from output to input using a more general mechanism calling pinning, denoted here by the red † next to $o_1 = \langle \text{Edge}, 38.5 \rangle$. A pinned output row intuitively narrows the entire execution down to only parts that are “relevant” to it (which we define formally later in Section 3). As shown in Figure 2, relevant rows in upstream tables are automatically shown in red. Specifically, input rows s_2 and s_3 join with f_1 and f_3 in a nested-loop fashion to produce rows j_2 through j_5 in the joined & filtered table; they are then further grouped into g_1 in the group table before finally producing o_1 in the output.

As soon the user pins o_1 , OURSYS automatically identifies the relevant input combos and “positions” execution at the very first input combo in lexicographical order — this is how the input combo (s_2, f_1)

Bar	Beer	Price	Drinker	Bar	Times
Edge	Amstel	4	Ben	Edge	4
Edge	Amstel	4	Dan	Edge	3
Edge	Corona	1.5	Ben	Edge	4
Edge	Corona	1.5	Dan	Edge	3

Displaying page 2 out of 512 pages

Figure 4: A paginated table in OURSYS.

was activated in the first place in Figure 2. Starting with this input combo, the user can then step through all other input combos relevant to o_1 , or manually choose any specific combo to investigate. Throughout the entire process, OURSYS automatically updates the highlighting in all downstream tables as well as the state of any expression evaluation trees, in effect supporting forward tracing.

When examining the execution for (s_2, f_1) as shown in Example 2.2, the user notices that (s_2, f_1) contributes a value of 16 to the final sum. According to the filter expression tree, the price of Amstel is higher than \$2, but $\text{EXISTS}(Q_{\text{inner}})$ returns true, meaning that Ben should like Amstel per query intention. OURSYS allows the user to “drill down” and explore the execution of Q_{inner} in this context. Recall that Q_{inner} is a correlated subquery. To an average programmer, the analogy is to think of evaluating Q_{inner} as a function call with an extra parameter setting for $f.drinker$, which takes its value from the current input combo. Hence, “drilling down” naturally corresponds to “stepping into” a function call in GPL debugging.

Once the user drills down to Q_{inner} , OURSYS creates a new panel for debugging this subquery block, illustrated in Figure 3. The UI makes it clear that we are executing

```
|| SELECT * FROM Likes l WHERE f.drinker = l.drinker;
```

with our parameters (i.e. $f.drinker$) set to Ben. In this case, a quick glance at the output table or the filter expression tree for this block should reveal the problem — nothing about this subquery supports that Ben likes Amstel. In fact, this subquery does not even know we are looking for Amstel. Therefore, to fix the query, we should let the outer query “call” Q_{inner} with an additional parameter (i.e. $s.beer$), and let Q_{inner} additionally test $s.beer = l.beer$.

The user can modify the original query accordingly and restart the debugging session to verify that it fixes the problem. Because the new query is syntactically similar to the old, OURSYS will produce a very consistent experience for the user: the execution of the new query will be nearly identical to the old, with the exact same stages and exact same ordering of input combos and intermediate result rows.

A Note on Scalability. While Example 2.2 simplifies our discussion by assuming a small database instance, realistic databases are much bigger. Even with moderately sized databases, joins can easily produce very large intermediate results. As we will see in Section 5, for some TPC-H benchmark queries, even with a moderate scaling factor of 1, a single intermediate result table can easily take an hour to print out entirely on the database server console. Hence, caching entire results at any location (database server, middleware, or user browser) or shipping them is simply impractical.

At the same time, users are generally unable to examine many rows simultaneously anyway. Therefore, OURSYS UI supports a pagination mechanism to display each table (input, intermediate

result, or final output); a screenshot is shown in Figure 4. The user only sees a page’s worth of data at once in each table. Data outside the visible page can be computed and fetched on demand (and subsequently cached or evicted). As the user interacts with the UI, OURSys adjusts the visible portion of all displayed tables accordingly, such that each reflects the execution point defined by the current input combination (combo). Similarly, the user can visit any pages of input tables to select a new input combo for forward tracing; in that case, OURSys would automatically adjust the displays of downstream tables to show the corresponding intermediate result rows.

Hence, efficient pagination is key to scalable SQL debugging. With efficient pagination, OURSys effectively allows the user to “teleport” across points of interest without incurring the execution cost in between, which makes OURSys much more powerful than GPL debugging. Pagination also provides a more manageable and less overwhelming experience for users. While most database systems support efficient pagination of input tables, it is challenging to paginate intermediate results as well as associated information for debugging. Furthermore, OURSys’s requirement of making execution reproducible imposes specific ordering of intermediate results that complicates optimization. We discuss how to tackle these challenges in Section 4.

3 DEBUGGING PARADIGM

This section describes the debugging paradigm of OURSys in detail. Several concrete instances of the models described are available in the full version [33]. OURSys supports a rich set of SQL query features, such as SELECT queries with inner cross joins, set/bag operations (e.g., UNION, INTERSECT, EXCEPT), and even outer joins and joins expressed in general JOIN syntax (except LATERAL). We also support both scalar and correlated subqueries. OURSys, however, currently does not support recursive WITH, LATERAL joins, WINDOW functions, and any built-in functions whose results cannot be reliably reproduced (e.g., RAND).

3.1 Data and Execution Model

3.1.1 Table Model and IIDs. OURSys allows users to interact with two types of tables: *base* tables and *derived* tables. Base tables are those that exist in the database, while derived tables are computed from the base tables during execution. To support reproducible execution (including ordering of intermediate result rows), we depart from the default unordered multiset semantics of SQL and instead model each table as an ordered list of rows, each associated with an *internal row id* (IID). A table’s IIDs must be drawn from a totally ordered domain and uniquely identify the rows within the table. Moreover, the IIDs do not influence the semantics of SQL query operators and should not be considered as extra columns by these operators. For example, if two rows have identical values for all columns (ignoring IID), they should be considered as duplicates even though their IIDs differ.

For a base table with a primary key declaration, we simply define IID to be its primary key value. Otherwise, we choose a compact UNIQUE key if one is available. In the worst case, if the table has no keys we use the database system’s internal row id (e.g., PostgreSQL’s `ctid`); such ids are unique even among duplicate rows.

For a derived table, we define its IID according to how it is computed. For each SQL query operator, we define a *canonical execution procedure* and a *result IID synthesis function* (more details will be provided shortly). The IID synthesis function computes the IID for each result row on the fly during canonical execution, such that the result rows are always produced in the IID order.

In addition to maintaining a reproducible order, OURSys uses IIDs in supporting a variety of debugging features. Therefore, good IID designs positively impact performance, as we will see later in Section 4. While it is technically possible to simply make the IID of a row its sequence number in the result set, such numbers by themselves do not provide any information useful for tracing or possible query rewrite optimizations. As we will see below, most of the IIDs in OURSys are “logical” instead of physical.

3.1.2 Query Blocks and Debugging Contexts. Given a query, OURSys defines a *canonical execution procedure* that is always followed when debugging. A complex query can be viewed as a collection of syntactic blocks with dependencies among them. OURSys models the canonical execution of such queries in terms of function calls –

- The outermost query block defines a function that computes the query result when executed.
- A subquery block defines a function that can be called by the function corresponding to its enclosing query block. A correlated subquery is analogous to a function with arguments.
- Each table defined by WITH is a function that computes the table contents when the table is referenced.

Overall, the canonical execution starts by executing the function defined by the outermost query block, which then calls functions corresponding to its subqueries or tables defined by WITH, which may further call functions for their subqueries, etc.

As a function may be called multiple times, for each *invocation* of a function, OURSys creates a *debugging context* as needed, analogous to an “activation frame” in a GPL. This debugging context holds information specific to the particular execution of the query block, such as the values of external column references.

3.1.3 Canonical Execution for Each Query Block. Having discussed the overall canonical execution procedure, we now zoom in on the canonical execution of each query block. Due to space constraints, we will only discuss SELECT blocks with inner cross joins. Because of the complexity of SELECT (including grouping and aggregation), we further decompose its execution into *stages*, where each stage can be seen as a query operator with its own canonical execution procedure and result IID synthesis function.

The first stage in a SELECT block is *join & filter*. Its canonical execution is nested for-loops iterating through rows, one for each input table in FROM, in order. In the innermost loop, we test the WHERE condition (which may involve calling subqueries, potentially with argument values obtained from the loop variables). The result row IID is synthesized as a vector whose components are the IIDs of the joining input rows. Note that the lexicographic order of these vector IIDs is consistent with the row production order.

If the block contains any grouping or aggregation, a *grouping* stage will be next. Its canonical execution is a stable sort of its input

rows according to the list of GROUP BY expressions¹ in order. Each result row starts with the GROUP BY expression values as columns, followed by additional columns needed to evaluate the remainder of the query (e.g. HAVING or SELECT expressions). The result row IID is synthesized as a vector whose components are the GROUP BY expressions in order, followed by the IID of the input row. Note that this order puts member rows of a group together, allowing the UI to detect group boundaries and display them specially (Figure 2). The *stable* sort ensures that the order of these vector IIDs is consistent with the row production order.

The final stage of the SELECT block produces the *final output* for the entire block. The canonical execution processes the input rows in order. If HAVING is present, input rows whose group does not pass the HAVING condition will be filtered out. Then, if the previous stage is a grouping stage, we produce one result row for each group of input rows, using the leading portion of the IID corresponding to GROUP BY expressions as the result IID.

3.2 Debugging Operations

We now describe what debugging operations a user can perform with OURSys, focusing on those requiring formalization and more in-depth discussion; others mentioned in earlier sections with straightforward semantics (such as visualization of expression tree evaluation) are omitted. Again, we will describe the operations mostly in the context of SELECT blocks with inner cross joins, although we have also generalized them to other SQL constructs.

3.2.1 Input Combination ("combo") Space and Execution Positioning. For each debugging context, we define an *input combo space* as an ordered set of *input combinations*, each representing a particular point in execution. The (conceptual) current point of execution is called the *active input combo* for the debugging context. For a SELECT debugging context with n input tables, the active input combo is the n input rows being examined inside the innermost loop by the canonical execution of the join & filter stage, and it is represented by an n -dimensional vector whose components are the IIDs of these input rows. For instance, the active input combo of the SELECT debugging context shown in Figure 2 is $\langle s_2, f_1 \rangle$, drawn from the input combo space $\{s_0, \dots, s_5\} \times \{f_0, \dots, f_3\}$.

The user can position the execution of a debugging context at a particular point using either *stepping* or *teleporting*. With *stepping*, OURSys automatically advances the active input combo to its successor (or predecessor if stepping in reverse order) in the input combo space. With *teleporting*, given any input table, through the paginated display illustrated in Figure 4, the user can jump or scroll to any page and select a particular row as active; OURSys will then update the active input combo accordingly.

3.2.2 Forward Tracing. An active input combo in the debugging context can produce *derivative* rows in the downstream result tables produced by the stages. Informally, the input combo contributes to the computation of derivative rows. As discussed in Section 2, once the active input combo is set, OURSys automatically refreshes all visualizations of expression trees, such that they reflect evaluation over the active input combo or its derivative rows. OURSys also

automatically refreshes all result tables in the debugging context, such that they show the pages containing and highlighting the derivative rows. This *forward tracing* feature allows the user to examine the effect of input rows on subsequent processing and potentially understand why the desired effect is not achieved.

Note that for a SELECT block, a given input combo can contribute to at most one result row per stage, so there is no ambiguity in which derivative rows to show downstream. By design, OURSys ensures this rule of at most one derivative per stage for other blocks (such as set/bag operations) as well.

3.2.3 Pinning: Tracing and Watchpointing. We first describe the semantics of *pinning* and then discuss its use for tracing and watchpointing. Consider all tables displayed for a debugging context. OURSys allows the user to *pin* up to one row from each of these tables. Formally, each pinned row defines a subset of the input combo space, called the row's *pinned (input combo) space*. Overall, the pinned space for the debugging context is its input combo space *intersected* with each of the pinned spaces defined by the pinned rows. Intuitively, the pinned space allows the user to narrow the execution down to the points of interest.

Consider a SELECT block joining n tables in FROM with input combo space $\mathbf{R} = \prod_{i=1}^n R_i$, where each R_i denotes the ordered list of IIDs for the i -th input table. A pinned row in the j -th input table with IID x defines a pinned space of $\prod_{i=1}^{j-1} R_i \times \{x\} \times \prod_{i=j+1}^n R_i$; i.e., the user is interested only in input combos with x participating. A pinned row in a result (intermediate or final) table defines a pinned space of $\{v \mid v \in \mathbf{R} \wedge x \text{ is a derivative row of } v\}$; i.e., the user is interested only in input combos that contribute to x .

Pinning has multiple uses. First, it augments OURSys's tracing capability: besides forward tracing from the active input combo, pinning effectively allows *backward tracing* from a pinned result row produced by any stage to the pinned space of input combos, and then from there, forward tracing to result rows further downstream. Second, combined with stepping, pinning provides a form of *watchpointing*. With a pinned space in effect for the debugging context, OURSys restricts stepping to the pinned space, effectively setting a watchpoint that pauses execution only at points relevant to the pinned rows.

3.2.4 Drilling Down and Pulling Up. As illustrated in Example 2.2, OURSys allows the user to *drill down* into a subquery, analogous to "stepping into" a function call. Besides drilling down into subqueries in WHERE, HAVING, and SELECT expressions, OURSys also allows drilling down into subqueries in FROM, which can be either directly nested therein or via a reference to some WITH definition. In these cases, the user can drill down directly through a particular row in a derived input table; OURSys will open the debugging context for the subquery responsible for producing that table and automatically pin that row in the subquery's final output table.

When debugging a complex query with many nested blocks, OURSys essentially maintains a "call stack" of debugging contexts. To let the user *pull up* from a subquery debugging context, OURSys simply returns the user to the previous debugging context on the stack, which belongs to the enclosing query block. The state of the subquery debugging context is still preserved until the user changes the active input combo in the enclosing block's debugging context, which forces "stepping out" from the last subquery function call.

¹An aggregate query without GROUP BY can be regarded as having an empty list of GROUP BY expressions.




Figure 5: OURSys Architecture

4 SYSTEM AND OPTIMIZATIONS

In this section, we describe the system implementation and optimizations for OURSys. Despite the myriad of details, there are three important high-level ideas. 1) Rather than showing the entire canonical execution of a query Q being debugged, it suffices to let the user examine one small, relevant window of this execution at a time. 2) To obtain all debugging information needed for a particular execution window, instead of performing the canonical execution and instrumenting it, we can formulate SQL queries based on Q to compute such information directly and declaratively. 3) We can judiciously compute some summary data and then use them to further rewrite these queries to be more efficient, without requiring special indexing support in the database itself.

Figure 5 shows the architecture of OURSys, which follows a client-server setup for easy deployment and better maintainability. The client side is a web frontend responsible for rendering and displaying data. It contains two primary interfaces: a query editor and a debugger interface. The server side is an API server in charge of computing data to support the debugger on the client side. Its debugging context generator parses all query blocks and initializes all debugging contexts in a query, and the data retriever responds to subsequent data requests from the debugger. Both services query the user-specified database to compute their corresponding responses to the client. In addition, the API server does not cache or maintain any front-end state in the database, making OURSys a plug-and-play tool that significantly improves usability.

4.1 Optimization Challenges

While it is possible to fetch complete debugging information (e.g., table content, data lineage, etc.) for the entire canonical execution of all query blocks through the debugging context generator at the initialization of their debugging context, such an approach is not scalable.

Example 4.1. Consider a TPC-H [10] database instance generated with a scale factor of 1 (i.e., data size of all tables is 1GB), and the following instantiated inner query of Q8 in the TPC-H benchmark:

```
|| SELECT EXTRACT(year from o_orderdate) as o_year,
||       l_extendedprice * (1 - l_discount) as volume,
||       n2.n_name as nation
|| FROM part, supplier, lineitem, orders,
||       customer, nation n1, nation n2, region
|| WHERE p_partkey = l_partkey AND s_suppkey = l_suppkey
|| AND l_orderkey = o_orderkey AND o_custkey = c_custkey
|| AND ... -- omit for simplicity
```

The minimum debugging information required to be poured into the client side is as follows:

There are ten stage tables fetched from the database: eight input tables, a join & filter table, and an output table. The data lineage

Item	Size
All table contents for all execution stages	1825 MB
Tuple-level data lineage	88 MB

contains two parts: the lineage between input tables and the join & filter table and the lineage between the join & filter table and the output table. While most browser tabs use less than 1GB of memory, the combined data here approaches 2GB. Allocating 2GB to a single browser tab is inefficient, especially when users examine queries in multiple tabs, each running as a separate process.

To summarize, the true challenge for OURSys is to accommodate large amounts of data with limited memory while ensuring scalability and interactive experience for query debugging. By focusing on what a user is capable of exploring at once, we state the intuition for an on-demand data retrieval solution:

- For each stage table, we implement a pagination technique and only retrieve the pages of tuples that the users focus on.
- For data lineage and expression evaluations, we retrieve them on a per-tuple level upon users' request, minimizing the computation on the server side to ensure an interactive experience.
- All data retrieval is done by rewriting part of or the entire original query, thus maintaining no state or materialized view in the database.

On-demand Retrieval Overview. On-demand data retrieval is done in two phases. In the first phase, for all stage tables in all query blocks, the debugging context generator creates metadata for quickly retrieving a particular page of a stage table and passes the metadata to the client side. While this can trigger a wait for inefficient queries as OURSys collects metadata on every stage table, the metadata helps guarantee interactive speed in subsequent debugging operations. In section 5, we show that collecting metadata is actually significantly more efficient than caching the entire table.

In the second phase, when a user starts to explore a particular page in any stage table, the client-side debugger makes requests to the server-side data retriever based on the metadata to fetch data only within the specified page. The computation overhead is relatively small since the page size (i.e., the number of tuples in a page) is usually much smaller than the entire table. If the user performs a debugging operation, the client-side debugger also relies on the metadata to fetch the corresponding data lineage that supports the operation, and such computation overhead is also insignificant on the server side, as debugging operations primarily act on a small number of tuples. As a result, OURSys avoids lengthy interruptions during a debugging session.

Baseline Approach for Page Retrieval. Before diving into the details of data retrieval, we realize that the current, most ubiquitous way of pagination is to rely on OFFSET and LIMIT to specify the page of interest. While this approach may work in fetching the leading pages of a query, it has significant drawbacks:

- The execution time grows rapidly as the OFFSET becomes larger. If a user is interested in the trailing pages of a table, using OFFSET and LIMIT essentially requires the database to execute the entire query every time the user jumps to a different page.
- If an explicit order is defined, using OFFSET and LIMIT has to execute the entire query to sort and determine the correct page,

which can take a long time and significantly worsen the debugging experience. Due to the necessity of stable orders on tables, OURSys cannot use OFFSET and LIMIT for pagination.

Cache Mechanism. Building on top of the on-demand data retrieval, OURSys caches computed pages and lineages on the client side to guarantee smooth browsing of frequently visited pages and avoid repetitive computation of the same content. It maintains an LRU-style cache whose max size is limited to 500 MB. Furthermore, OURSys caches the results of all independent scalar queries (given the minimal size of their results) and rewrites the query to replace the scalar query with the cached constant at the time of execution to further speed up the query.

4.2 Efficient Data Retrieval

OURSys collects statistical summaries for each page (done by the debugging context generator) and later uses the summaries as potential pushdown predicates to efficiently fetch a specific page through the on-demand data retriever. The metadata of a stage table contains two parts: (1) an ordered list of *milestones*, each of which contains statistical summaries of a page, and (2) a *table query* that produces the entire output of the table.

The composition of a milestone is as follows:

- The sequence number of the first tuple in the page to keep track of the page offset in the entire table.
- Minimum tuple IID to locate the first tuple on the page.
- Tuple count of the page for the client side to know the page size of the last page in case the total tuple count is not divisible by the user-specified page size (i.e. in situations where the last page has fewer tuples than the other pages).
- The minimum and maximum of all columns with an index in the database on the page which later provides the query optimizer with hints for potential optimization when fetching a page.
- A bloom filter on the external bindings (i.e., the columns referenced in a correlated subquery) only if correlated subqueries exist. The bloom filter is implemented as an extension to the database, and the purpose is to “short-circuit” the query on the input combos that do not evaluate to true on the subqueries, thus avoiding expensive subqueries.

To obtain milestones, the debugging context generator creates and executes a *milestone query* for all tables. Then, the context generator pairs the milestones with a corresponding table query and ships them to the client. The milestone query of a table is composed as follows:

- (1) Given the table query, first replace the SELECT clause with a sequence number, a tuple IID, the number of tuples, and all indexed columns. Use the window function ROW_NUMBER() to create the sequence number and order the tuples by IID.
- (2) Given the above query result, group the result by the page size, then project the minimum sequence number, minimum IID, and all the statistics summaries.

Example 4.2. Consider the following TPC-H query:

```
|| SELECT p_name, COUNT(DISTINCT o_orderkey) num_order
|| FROM part, lineitem
|| WHERE p_partkey = l_partkey
|| AND EXISTS (
||   SELECT * FROM orders WHERE o_orderkey = l_orderkey)
|| GROUP BY p_name
```

```
|| ORDER BY p_name, num_order;
```

The following is an example of the join & filter table query, which evaluates the implicit join (in the FROM) and the conditions (in the WHERE):

```
|| -- table query
|| SELECT * FROM part, lineitem
|| WHERE p_partkey = l_partkey
|| AND EXISTS (SELECT * FROM orders WHERE o_orderkey = l_orderkey);
|| -- milestone query
|| WITH tmp(seq, iid, p_partkey, l_orderkey, l_linenum) AS (
||   -- modified table query
||   SELECT ROW_NUMBER() OVER (ORDER BY p_partkey, l_orderkey, l_linenum) - 1,
||         ROW(ROW(p_partkey), ROW(l_orderkey, l_linenum)),
||         p_partkey, l_orderkey, l_linenum -- indexed columns
||   FROM part, lineitem
||   WHERE p_partkey = l_partkey
||   AND EXISTS (SELECT * FROM orders WHERE o_orderkey = l_orderkey))
|| -- MIN_IID is a user-defined function to find the minimum IID
|| SELECT MIN(seq), MIN_IID(iid), COUNT(*) count,
||        MIN(p_partkey), MAX(p_partkey), -- statistical summaries
||        MIN(l_orderkey), MAX(l_orderkey),
||        MIN(l_linenum), MAX(l_linenum),
||        BLOOM_FILTER(l_orderkey)
|| FROM tmp
|| GROUP BY seq / 50 -- page size 50
|| ORDER BY seq / 50;
```

Suppose the milestone query returns the following results; headers are abbreviated:

seq	min IID	cnt	p_key	o_key	line #	bloom
0	((1), (1,3))	50	[1,9]	[1,4]	[3,6]	00100111
50	((9), (3,5))	20	[9,11]	[2,5]	[3,6]	10101010

When a user wants to retrieve a specific page, the client rewrites the table query to a *page query* by using the IIDs in the milestones to precisely define the page as well as appending extra pushdown conditions to the end of the table query based on the statistic summaries (i.e., min/max of the indexed columns). Once the rewrite is done, the client delivers the rewritten query to the on-demand data retriever, and the retriever executes the query and sends back the results. This process is also illustrated in Figure 5. Below is a sample query to fetch the first page of the join & filter table:

```
|| SELECT ROW(ROW(p_partkey), ROW(l_orderkey, l_linenum)), -- tuple IID
|| *
|| FROM part, lineitem
|| WHERE p_partkey = l_partkey
|| AND EXISTS (
||   SELECT * FROM orders
||   WHERE o_orderkey = l_orderkey
||   -- extra pushdown to subquery
||   AND o_orderkey BETWEEN 1 AND 4
||   -- IID guarantees a page of 50
||   AND ROW(ROW(p_partkey), ROW(l_orderkey, l_linenum)) >= ((1), (1,3))
||   AND ROW(ROW(p_partkey), ROW(l_orderkey, l_linenum)) < ((9), (3,5))
||   -- indexed column pushdown
||   AND p_partkey BETWEEN 1 AND 9
||   AND l_orderkey BETWEEN 1 AND 4
||   AND l_linenum BETWEEN 3 AND 6
||   AND BLOOM_TEST(00100111, l_orderkey) -- if false, no need to evaluate EXISTS
|| ORDER BY p_partkey, l_orderkey, l_linenum;
```

Lineage Retrieval. When fetching a page, the page query also retrieves the lineage of each tuple in the page. In the previous example, there is no explicit SELECT expression that projects the lineage because the IID of the join & filter table is the lineage, as it is equivalent to the input IID combo. However, it is not clear how to retrieve the lineage of a group in the group table since a group consists of multiple joined tuples, and it is infeasible to aggregate all joined tuples as the lineage since the data size can be large. OURSys

chooses the minimum joined tuple IID as the lineage, and we show how this choice helps with pinning later.

Remark. OURSys performs recursive predicate pushdown for subqueries. As shown in the previous example, it is natural to push down `l_orderkey` to the subquery. Furthermore, OURSys also pushes down predicates to derived input tables, i.e., OURSys only fetches the necessary input pages from the derived input table when computing downstream stage tables in the outer query. However, such a pushdown mechanism does not work when the range of the indexed column is too large or when the false positive rate of the bloom filter is too high. Therefore, OURSys sets a threshold for both range overlap (default to 0.3) and the false positive rate (default to 0.5). If the overlap or false positive rate is higher than their corresponding threshold, the corresponding condition will be removed to prevent ineffective predicate pushdown that can potentially slow down query execution.

4.3 Debugging Support

We now describe how OURSys supports in-context debugging operations (e.g., forward tracing, pinning, stepping).

4.3.1 Forward Tracing. Forward tracing enables users to select a particular input combo from each input page and observe how it propagates through the execution stages. Therefore, we safely assume that the client side has already fetched the input pages that contain the input combo before users can move the cursor of each input table to the input tuples of interest.

Knowing the input combo, OURSys first computes the downstream tuples (i.e., tuples in the join & filter, group, and output table) this input combo contributes to, and then it locates the pages that contain the downstream tuple. With the located pages, the client-side debugger can then request the page through the data retriever following the previous procedures. Thus, forward tracing is achieved in two phases. In the first phase, the client sends the selected input combo along with all downstream table queries to the data retriever, which then rewrites the table queries to compute the corresponding downstream tuples in the downstream stage tables and ships them back to the client. The query rewrite is done in the following order:

- (1) The input IID combo is formulated as a condition and appended to the join & filter query. The data retriever executes the query. If the query returns an empty result, we stop here, and no other query needs to be executed as this input combo does not satisfy the WHERE clause/join condition.
- (2) If the previous query returns non-empty results, the input combo must contribute to a group. Therefore, the input IID combo can be injected into the group table query. By executing the query, we obtain the group's IID.
- (3) Inject the group's IID as a WHERE condition to the output table query. By executing the rewritten output query, we can compute the IID of the output tuple. The reason why we use the group's IID instead of the input IID combo is that the group formed by a single input combo might not pass the HAVING clause.

Consider the query in Example 4.2. Suppose the selected input combo is $((1), (1, 3))$, and it contributes to the group ("Nail").

The on-demand data retriever rewrites the queries for the first phase as follows:

```
-- join & filter
SELECT ROW(ROW(p_partkey), ROW(l_orderkey, l_linenum)), -- tuple IID
       p_partkey, l_partkey, p_partkey = l_partkey, -- expression evaluation
       EXISTS(SELECT * FROM orders WHERE o_orderkey = l_orderkey),
       p_partkey = l_partkey AND EXISTS(SELECT * FROM orders WHERE o_orderkey
                                       = l_orderkey)
FROM part, lineitem
WHERE p_partkey = l_partkey
AND EXISTS(SELECT * FROM orders WHERE o_orderkey = l_orderkey)
ROW(ROW(p_partkey), ROW(l_orderkey, l_linenum)) = ((1), (1,3));
-- group
SELECT ROW(p_name) -- tuple IID
FROM part, lineitem
WHERE ROW(ROW(p_partkey), ROW(l_orderkey, l_linenum)) = ((1), (1,3))
GROUP BY p_name;
-- output table
SELECT ROW(p_name) -- tuple IID
FROM part, lineitem
WHERE p_partkey = l_partkey
AND EXISTS(SELECT * FROM orders WHERE o_orderkey = l_orderkey)
AND ROW(p_name) = ROW('Nail')
GROUP BY p_name;
```

Note that here OURSys intentionally augments the join & filter query further to include the evaluation expression of all nodes on the predicate tree, so that this information can also be sent to the client side to render the predicate tree as shown in Figure 2.

In the second phase, the client uses the computed downstream tuples to locate the page (since tuples are sorted by IID, this can be quickly done via binary search) and requests those pages through the data retriever, which then efficiently fetches the pages based on the mechanism described previously.

4.3.2 Pinning. In the case of *SPJ* queries, pinning is essentially the same as forward tracing, as pinning any tuple in the join & filter table or output table traces back to only one input combo. Since the IIDs of these two tables are the input IID combos, the client side first fetches the corresponding input pages by decomposing the IID of the pinned tuple and then invokes the procedure of forward tracing.

For *SPJA* queries, pinning a tuple in the join & filter table is the same as forward tracing, as it traces back to only one input combo. Pinning in a group table or an output table automatically forward traces the first input combo in the pinned input combo space, and such information is returned by the page query as shown in Example 4.2. However, additional computation is required to compute if a tuple in the page is in the pin space, as the pin space for each table potentially contains more tuples. OURSys achieves such computation by inserting an extra Boolean SELECT expression in the page query as the flag to indicate if a tuple is in the pin space.

The inserted Boolean SELECT expressions differ for different stages, and OURSys works backward when determining the pin space for each table. Assuming a pin is placed in the output table, we can then use its IID to determine the corresponding group in the group table and create a corresponding equality condition as the flag SELECT expression. With the evaluation of the GROUP BY expressions, we can further create equality conditions as flags for the join & filter table to check if a single input combo contributes to the designated group. Finally, we check if an input tuple is in the pin space by verifying its contribution to the pin space of the join & filter table.

Consider the query in Example 4.2 again. The data retriever rewrites the page queries for upstream tables as follows, assuming a pin is placed on the output tuple ("Nail", 30):

```
-- group
SELECT ROW(p_partkey), p_name,
       p_name = 'Nail' -- Boolean flag indicates pin space
FROM part p1, lineitem l1
WHERE ROW(p_name) >= 'Hammer'
AND ROW(p_name) < 'Screwdriver' -- define the page, 'Nail' is included
AND ... -- original WHERE predicates and pushdown conditions
GROUP BY p_name;

-- join & filter
SELECT ROW(p_partkey), *,
       p_name = 'Nail' -- Boolean flag indicates pin space
FROM part, lineitem
WHERE ROW(ROW(p_partkey), ROW(l_orderkey, l_linenum)) >= ((1), (1,3))
AND ROW(ROW(p_partkey), ROW(l_orderkey, l_linenum)) >= ((9), (3,5))
AND ... -- original WHERE predicates and pushdown conditions

-- input part, lineitem follows a similar style
SELECT ROW(p_partkey), *,
       EXISTS(
         SELECT * FROM part p2, lineitem l2
         WHERE p2.p_partkey = l2.l_partkey
         AND EXISTS (SELECT * FROM orders WHERE o_orderkey = l2.l_orderkey)
         AND p_name = 'Nail'
         AND ROW(p1.p_partkey) = ROW(p2.p_partkey) -- match IID
       ) -- Boolean flag indicates pin space
FROM part p1
WHERE ROW(p_partkey) >= ROW(1) AND ROW(p_partkey) < ROW(9);
```

As shown above, the extra SELECT expressions are inserted into the page queries of the corresponding tables. Therefore, it causes only a small extra computation overhead and does not affect the overall interactive experience of OurSys.

4.3.3 Stepping. Stepping can be performed in only two scenarios: (1) when there is no pin in any table, users essentially step forward/backward through the entire input combo space in the nest-loop order, (2) when there are pins in tables but the pinned input combo space has not been restricted to only one input combo (e.g., pinning a group).

Stepping through the entire input combo space can easily be achieved by fetching new pages followed by forward tracing because the next input combo is predictable, given the nest-loop order. On the contrary, stepping through a pinned input combo space needs special accommodation as the following input combo in the pinned input combo space cannot be predicted.

Therefore, the idea is to step in a pinned input combo space by using the pinned tuples as extra predicates to compute the following input combo given the current input combo. Since the IID of the join & filter table represents the input combo, the problem can be converted to find the next joined tuple in the pin space of the join & filter table, and the rewrites follow naturally.

Consider the query in Example 4.2 again. The data retriever rewrites the query of the join & filter table, assuming a pin is placed on the output tuple ("Nail", 30):

```
SELECT ROW(ROW(p_partkey), ROW(l_orderkey, l_linenum))
FROM part, lineitem
WHERE p_partkey = l_partkey
AND EXISTS (SELECT * FROM orders WHERE o_orderkey = l_orderkey)
AND ROW(ROW(p_partkey), ROW(l_orderkey, l_linenum)) >
  %(current_input_combo_id)
AND p_name = 'Nail' -- pin-induced predicate
ORDER BY p_partkey, l_orderkey, l_linenum
LIMIT 1;
```

The above query computes the next input combo in the pinned input combo space, and the client can invoke the forward tracing procedure to move the input cursors to the next input combo. If the users usually take multiple steps in the input combo space, the client can choose to increase the LIMIT and cache the results.

5 PERFORMANCE EXPERIMENTS

We conducted experiments to evaluate the performance of our optimizations in Section 4. To compare our optimization with the baseline approach (i.e., OFFSET and LIMIT) on the efficiency of fetching pages, we ran three evaluations independently to examine how much performance improvement can be achieved

- via bloom filter by short-circuiting correlated subqueries.
- via predicate pushdown based on the statistical summaries.
- in general by combining all the optimizations (i.e., predicate pushdown, bloom filter, and cached independent scalar query).

Workload and Tasks. To evaluate the above items independently, we used TPC-H [10] as our benchmark database. We generated three database instances of sizes 1GB, 5GB, and 10GB. In addition to the primary indexes on the primary keys of each table, we also created a reasonable set of secondary indexes across multiple tables in the instances (details in Appendix A) to reflect a more realistic scenario. For the first two evaluations, we cherry-picked one query for each to show the effectiveness of the corresponding optimization. For the general evaluation, we applied all possible optimizations to all 22 queries and reported general results.

Test Environment. All experiments were done locally on a 64-bit Ubuntu 22.04 LTS server with Intel(R) Xeon(R) Gold 6138 CPU @ 2.00GHz, 64 GB RAM, and 256 GB disk space. We used PostgreSQL-16 [54] and only changed the following configurations: We turned off parallelism for a fair comparison, set the work_mem to 128MB, and set the shared_buffers to 8GB.

5.1 Bloom Filter Evaluation

We chose the join & filter table query from Q2 in the TPC-H benchmark for bloom filter evaluation:

```
SELECT ROW(p_partkey, s_suppkey, ps_partkey, ps_suppkey,
          n_nationkey, r_regionkey), *
FROM part, supplier, partsupp, nation, region
WHERE ... AND ps_supplycost = (
  SELECT MIN(ps_supplycost)
  FROM partsupp, supplier, nation, region
  WHERE p_partkey = ps_partkey -- p_partkey is external binding
  AND ...)
```

p_partkey is the only external binding we created a bloom filter for in the milestone query, which we applied in the page query to quickly fetch the page. By identifying if a specific p_partkey is in the page of interest, the page query is able to potentially avoid evaluating the expensive correlated subquery since the WHERE clause of Q2 is conjunctive. For the bloom filter, we set the number of bits (denoted by m) to 1024 by default, the estimated number of unique p_partkey in each page (denoted by n) was set to the page size, and the number of hash functions needed (denoted by k) was calculated by the formula: $k = \frac{m}{n} \times \ln 2$.

We ran two sets of experiments. First, we fixed the database size to 1GB and varied the page size. Second, we fixed the page size to 50 and varied the database size.

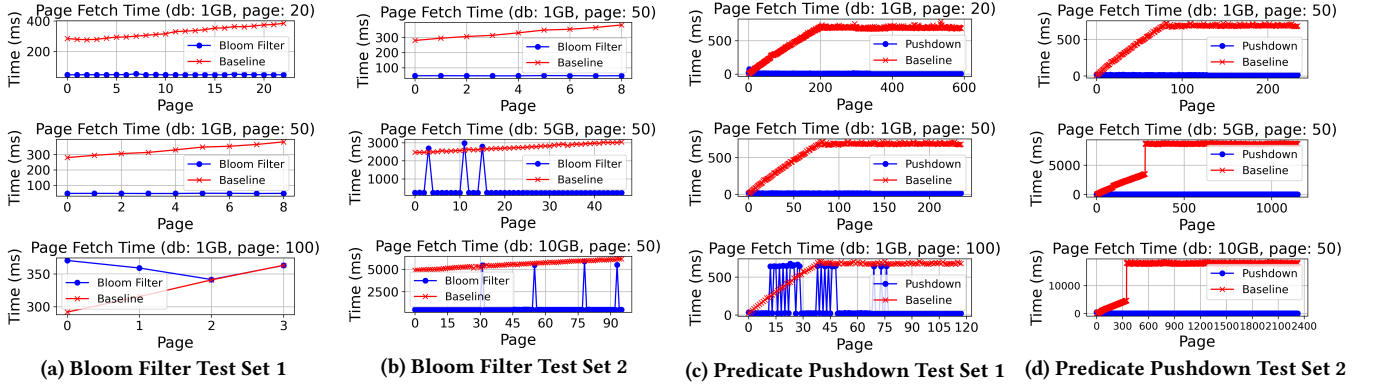


Figure 6: Experiments for Independent Evaluation of Bloom Filter and Predicate Pushdown

In both sets of experiments, we collected the execution time reported by the EXPLAIN ANALYZE command for all pages and observed its trend. We compared the bloom filter optimization with the baseline approach which use OFFSET and LIMIT to fetch a page. The results for both sets of experiments are shown in Figure 6a and Figure 6b, from which we derive the following conclusions:

- When bloom filter is effective, it outperforms the baseline approach. For page size of 50, the bloom filter keeps the execution time of the page fetch at around a tenth of the execution time of the baseline.
- Bloom filter is sensitive to page size as shown in Figure 6a. As the page becomes larger, it is likely more unique values are inserted into it, increasing the false positive rate. In the case where the bloom filter is ineffective, the DBMS pays extra cost to evaluate it with no benefit returned, causing the bloom filter to run more slowly with a page size of 100.
- Bloom filter is rather insensitive to the database size as shown in Figure 6b. While there are a few outliers, bloom filter performs fairly stably overall.

5.2 Predicate Pushdown Evaluation

We chose the join & filter table query from Q7 in the TPC-H benchmark for evaluation of predicate pushdown:

```

|| SELECT ROW(s_suppkey, l_orderkey, l_linenum, o_orderkey,
||           c_custkey, n1.n_nationkey, n2.n_nationkey), *
|| FROM supplier, lineitem, orders,
||      customer, nation n1, nation n2
|| WHERE ...

```

In addition to the primary keys of each table, we created pushdown conditions for any secondary indexed columns. Before executing the page query, we checked the pushdown range of the columns and removed all conditions whose ranges were more than 30% of the active domain of the corresponding columns, and this helped prevent us from “over-hinting” on the database and causing performance to degrade subsequently. We ran two sets of experiments similar to the evaluations for the bloom filter. The results are shown in Figure 6c and Figure 6d, based on which we can have the following conclusions:

- While predicate pushdown might perform similarly as the baseline for fetching the leading pages, its execution time does not grow linearly, resulting in a huge advantage over the baseline

when fetching middle pages or trailing pages in a table. As shown in Figure 6c and Figure 6d, such an advantage can save about two orders of magnitude on the execution time as shown in Figure 6c when the database size is 10GB and the page size is 50. The execution times for the baseline and pushdown optimization are around 42000ms and 500ms for trailing pages, respectively.

- Based on Figure 6c, predicate pushdown is also sensitive to the page size. As larger pages contain larger ranges, it becomes less likely to trigger a fast index scan in the database.
- Based on Figure 6d, the predicate pushdown optimization is not sensitive to the database size, and it provides great performance even if the database size is fairly large.

5.3 General Evaluation

For the general evaluation, we applied the optimizations to all queries in the TPC-H benchmark. The predicate pushdown optimization can be applied to all queries, while bloom filter and scalar caching can be only applied to Q4, Q16, Q18, Q20-22 and Q2, Q15, Q11, Q22, respectively. For all tables in all the query blocks of each TPC-H query, we collected the following data:

- The execution time of the milestone query, and size of its output.
- The execution time of the table query itself, and size of its output.
- The execution time of page query and the baseline query (i.e., OFFSET and LIMIT) for the first page, middle page and the second last page of the table. We chose the second last page instead of the last page to guarantee the fetch of a full page for fair comparison.

5.3.1 Milestone Query. Due to the space constraints, Table 1 shows a partial comparison between top-3 longest-running milestone/table query pairs with 1GB database and page size of 50 (complete results in [33]). All of them are from Q18, which is the most expensive query in the benchmark.

Q18	Milestone Query		Table Query	
	Exec. Time (ms)	Output (MB)	Exec. Time (ms)	Output (MB)
join & filter	53581.53	52	17478.41	1272
group	19449.81	9	18503.99	70
output	20030.26	9	18415.74	75

Table 1: Comparison Between Milestone and Table Query

In summary, the milestone query usually runs longer than the table query as it performs extra computation on top of the table query. However, its output is usually much smaller than that of the table, making it easier to process for the client side.

Q18 join & filter	1GB		5GB		10GB	
	Opt.	Base	Opt.	Base	Opt.	Base
head page	24.21	3876.84	255.36	34380.39	10617.46	158664.23
middle page	25.32	6863.59	326.19	72862.8	9497.6	181742.56
tail page	19.42	9850.34	244.22	111345.2	9660.1	204820.89

Table 2: Page Fetch Time Comparison between OURSys and Baseline. Unit: millisecond

5.3.2 Page Fetching for All Queries. To show the scalability of OURSys we show the execution time comparison between the page query and baseline query for Q18 (the most expensive query in the benchmark) across all database instances in Table 2, with a fixed page size of 50. While the baseline can potentially take minutes to return a page on a large database, OURSys returns the results in at most 10 seconds. The full results for all queries are presented in the Appendix A. In summary, the conclusion follows from the independent evaluation of bloom filter and predicate pushdown: the combined optimization is generally sensitive to page size and insensitive to the database size. It performs no worse (sometimes better) than the baseline approach when fetching leading pages, and it always significantly outperforms the baseline approach when fetching a middle/trailing page.

Final Remark. The execution times reported in the experiment are taken from the EXPLAIN ANALYZE command, which only measures the time spent by the PostgreSQL process. Such execution times do not include the time between sending the command and receiving results, which accounts for network latency. We intentionally chose to exclude such time from our experimentation because the output can potentially be quite large and a test conducted took an hour for the psql client to fully receive the entire output of the join & filter table of Q18. This is yet another strong indication that executing the table query and caching the entire output is inadequate and entirely unsuitable for interactive debugging.

6 USER STUDY

We conducted a user study in an undergraduate database course to evaluate the overall effectiveness of OURSys on two aspects: (1) does OURSys help users catch more logical bugs? (2) does OURSys reduce the time to identify logical bugs?

Participants. We recruited 237 students from the course, who were SQL beginners and had just become familiar with SQL at the time of the user study. The participation was voluntary except for the incentive of extra exercise to practice debugging skills. We considered the possibility of recruiting participants from other sources (e.g. Amazon Mechanical Turk) but decided against it because it was hard to quantify and control participants’ SQL familiarity to a similar level, potentially yielding inaccurate results on OURSys as SQL familiarity has a significant impact on the debugging time.

Preparation and Setup. The user study was conducted during two consecutive weekly 75-minute discussion sessions. In the first session, students were given a tutorial on OURSys and informed about the format of the survey-style quiz, which contained two SQL debugging questions. After the first discussion, OURSys is made public for students to get familiar with. In the second session, students completed the quiz synchronously in a proctored environment, and they were asked not to discuss the questions with classmates. For each question in the quiz, students were provided a problem statement, an incorrect query, its incorrect output, and the expected correct output. Q1 featured UNION and aggregation, containing two




Figure 7: Bugs Caught Distribution for Q1




Figure 8: Bugs Caught Distribution for Q2

logical bugs. Q2 featured OUTER JOIN, NULL, GROUP BY, and EXISTS, containing three logical bugs. The testing database instance in the user study had been made public and installed by students early in the course. The incorrect and expected outputs were generated based on that instance. We also prepared a remote pgAdmin web interface for students during the debugging sessions to free them from setting up their own instances, eliminating potential noise in the measurement of debugging time.

Tasks. To create treatment and control groups, students received the two questions in a random order. For the first question they received, they were free to use any tool of their choice (e.g., Gradescope autograder, Postgres console, pgAdmin web interface) except for OURSys. For the second question, OURSys was made available along with all other tools, and students had the option to use/not use OURSys and were asked to indicate it in the survey. For each question, students were to identify the mistakes in a free response text box. Students self paced, but they were recommended to spend about 15-20 minutes on each question. Students were not allowed to move onto the second question until they answered the first.

Results and Analysis. We collected data on the time it took students to answer each question, whether they chose to use OURSys and the mistakes they noted. Of 237 students, 140 participated in the study with complete and valid responses to both questions. Therefore, we based our analysis on these 140 responses. In summary, 73 students received Q1 first, and 67 received Q2 first, and they completed the question without OURSys. Consequently, 73 and 67 students received Q2 and Q1, respectively, as the second question with the OURSys option. When OURSys was available, 36 (for Q2) and 29 (for Q1) students chose not to use OURSys. Therefore, for Q1, we have 38 submissions with OURSys and 102 without OURSys (29 of them had OURSys option); for Q2, we have 37 submissions with OURSys and 103 without (36 of them had OURSys option).

The authors manually reviewed each valid response and verified the correctness of student-identified logical bugs. Given two bugs in Q1 and three in Q2, the distributions of identified bugs are shown in Figure 7 and Figure 8 respectively. Each chart shows how many

submissions caught the specified bugs for all possible combinations of caught bugs. While the numbers of submissions differ between using/not using OURSys, the distributions were similar. In fact, the average number of bugs identified are similar regardless of the usage of OURSys (shown in the last column of Table 3). Though this observation does not indicate OURSys helped students find more bugs, we notice a drastic drop in the debugging time as shown in Table 3, where using OURSys reduces the average debugging time for Q1 by 464.14 seconds and Q2 by 510.49 seconds. Based on the comparison between the bug combination distributions and the comparison between the debugging time, we conclude that using OURSys significantly improves debugging efficiency by helping users identify mistakes faster. The fact that all students had a similar familiarity with SQL further strengthens this argument.

Query	# Responses	Avg. Time (s)	Avg. Bugs
Q1 w/ OURSys	38	771.52	1.13 / 2
Q1 w/o OURSys	102	1235.66	1.18 / 2
Q2 w/ OURSys	37	1147.04	1.91 / 3
Q2 w/o OURSys	103	1657.53	1.85 / 3

Table 3: Average Time Taken and Bugs Caught per Question

7 RELATED WORK

Query Semantics Debugging. There are two main lines of work toward debugging query semantics. The first line helps users debug wrong queries against a correct or reference query to find the query syntax that cause semantic difference. Therefore, this line differs from OURSys fundamentally. XData [22] checks the correctness of a query by running the query on self-generated testing datasets. Cosette [24–26], SQLSolver [31] and QED [60] focuses on testing query equivalence using constraint solvers and theorem provers. RAtest [50] and Cinstance [34] aim at constructing small and illustrative database instances to show the differences between two queries. [21] developed a grading system that canonical-izes queries with rewrite rules and then decides query similarity based on a tree-edit distance between logical plans. SQLRepair [55] and QR-Hint [40] fix the wrong query by proposing syntax edits.

The second line, more related to OURSys, focuses on debugging a query without a reference query. Qex [59] is a tool for generating input relations and parameter values for unit-testing parameterized SQL queries. SQLLint [12–14] detects and alerts users about suspected semantic errors in a query, but it does not help users step through the execution. Habitat [29, 37] is a query execution visualizer that allows users to highlight parts of a query and view their intermediate results, but it does not provide any explanation on how and why some results are produced. DESQL [39] is a debugger for SQL in Spark, which decomposes the query into subqueries and help users examine the subqueries output, but it does not provide explanation on how outputs are produced. While I-Rex [41, 49] built an interactive interface with several debugging operations to step through execution, it does not accommodate large-scale data.

The following areas also contribute ideas to query debugging:

Data Exploration. DataPlay [5] lets users directly manipulate a graphical query by changing quantifiers and dependencies among constraints, and provides real-time updated query answers/non-answers to show how the output reacts to changes. As users input

various queries, AIDE [30] and SQL QueRIE [7] learns on users’ query history and “predicts” queries that retrieve interesting data for users. Explique [46] extends users’ queries by suggesting an extra set of possible selection predicates that potentially help users zoom into a specific area of the original answers. Toward the purpose of semantic debugging, Caballero et al. [16, 17] builds a computation tree which represents the execution flow of a query, guides users to annotate the output tuples and notified users when unexpected results are produced.

Data Provenance. Data provenance [11, 15, 23, 36, 42] provides a natural way to trace query execution. It gives explanation to why and how a particular output is or is not produced. Many frameworks have been built to capture provenance in relational systems [6, 27, 45]. More recently, the Perm [35] and GProM [9] system extends PostgreSQL with support for provenance capture implemented using instrumentation and relational encoding, respectively. PUG [47] extends GProM with support for capturing and summarizing provenance for queries with negation. Smoke [56] implements provenance capture in a main-memory database system. In addition, many provenance-capturing frameworks [8, 28, 43] target big data platforms like Spark and MapReduce.

Query Visualization. Many interactive query builders ([1–5]) employ block diagrams to show the interactions among tables and their attributes. This line of work focuses on helping users construct queries from scratch, and users can manipulate the diagrams to build desired queries. Starburst [38] develops the Query Graph Model to help users understand query plans, but the generated graphs are usually disconnected from the query syntax; GraphQL [20] and Visual SQL [44] follow the paradigm of entity-relationship representation to create query diagrams. QueryVis [48] creates unambiguous diagrams to capture comprehensive semantics of complex SQL queries, and has shown such diagrams effectively improve users’ ability to understand query semantics. SQLVis [51] shares similar approach with QueryVis, relying on graphs to express the relationship among tables and attributes.

While the above approaches have demonstrated significant advantages for some specific tasks, they are not comprehensive enough for SQL debugging in a general and organized way. Browsing through intermediate/final query outputs enables users to observe and reflect directly on their queries, but it lacks context on how the output is produced. While data provenance seems to be a good complement, it relies on a query plan to identify how and why tuples are produced or missing during query execution and does not directly map plan operators to the original query syntax, making it difficult to reveal problematic SQL constructs, especially for SQL novices. Both approaches are incapable of handling large amounts of data as 1) computing provenance for large data is infeasible, and 2) query optimizers might generate different plans for the same query, causing arbitrary and unstable order of the output rows that is not cognitively possible for users to keep track of. Query visualization usually requires users to learn new representations, which potentially creates extra cognitive burden.

Efficient Query Evaluation. Besides the debugging paradigm, OURSys also accommodates large data. However, most of the related work focuses on optimizing the query evaluation without

previous knowledge of the query, and this differs fundamentally from our setting. The line of work in direct access to query answer [18, 19, 32, 58] provides methods to access a random page in the query result, but it is restricted to conjunctive queries only. Therefore, it is not applicable in OURSYS as OURSYS needs to support large SQL fragments. Previous work on data skipping and predicate pushdown [52, 53, 57, 61] requires significant changes in the database systems. As a result, their proposed techniques are not applicable in OURSYS.

8 CONCLUSION AND FUTURE WORK

As SQL continues to be a leading querying and programming language, tools for understanding SQL queries maintain their importance both for helping novices understand and learn relational queries and to aid experts and practitioners in debugging their own complex queries. Building upon the work discussed in the previous sections, we present OURSYS, an interactive logical SQL debugger that allows users to examine the execution of SQL queries following the way they are written and debug them in a general and organized way. We define a debugging paradigm for SQL queries by making analogies to the debugging of general-purpose programming languages (GPLs). We develop multiple powerful debugging operations that support tuple-level examination of the execution. We designed and developed OURSYS to accommodate debugging over large databases while maintaining an interactive experience. We conducted performance experiments and a user study to demonstrate the efficiency and effectiveness of OURSYS. The development of OURSYS inspires multiple directions for future work, including but not limited to: (1) design and develop more debugging operations to better support debugging over large databases, and (2) support the examination of runtime errors, to which database systems usually provide uninformative hints (e.g. “ERROR: division by zero”).

REFERENCES

- [1] 2023. dbForge. <https://www.devart.com/dbforge/mysql/querybuilder/>.
- [2] 2023. Microsoft Access. <https://www.microsoft.com/en-us/microsoft-365/access>.
- [3] 2023. PgAdmin. <https://www.pgadmin.org/>.
- [4] 2023. Rapid SQL. <https://www.idera.com/rapid-sql-ide/>.
- [5] Azza Abouzied, Joseph Hellerstein, and Avi Silberschatz. 2012. Dataplay: interactive tweaking and example-driven correction of graphical database queries. In *Proceedings of the 25th annual ACM symposium on User interface software and technology*. 207–218.
- [6] Parag Agrawal, Omar Benjelloun, Anish Das Sarma, Chris Hayworth, Shubha Nabar, Tomoe Sugihara, and Jennifer Widom. 2006. Trio: A system for data, uncertainty, and lineage. In *Vldb*, Vol. 6. 1151–1154.
- [7] Javad Akbarnejad, Gloria Chatzopoulou, Magdalini Eirinaki, Suju Koshy, Sarika Mittal, Duc On, Neoklis Polyzotis, and Jothi S Vindhiya Varman. 2010. SQL QueRIE recommendations. *Proceedings of the VLDB Endowment* 3, 1-2 (2010), 1597–1600.
- [8] Yael Amsterdamer, Susan B Davidson, Daniel Deutch, Tova Milo, Julia Stoyanovich, and Val Tannen. 2011. Putting lipstick on pig: Enabling database-style workflow provenance. *arXiv preprint arXiv:1201.0231* (2011).
- [9] Bahareh Sadat Arab, Su Feng, Boris Glavic, Seokki Lee, Xing Niu, and Qitian Zeng. 2018. GProM-a swiss army knife for your provenance needs. *A Quarterly bulletin of the Computer Society of the IEEE Technical Committee on Data Engineering* 41, 1 (2018).
- [10] TPC Benchmark. [n.d.]. <http://www.tpc.org/tpch>.
- [11] Nicole Bidoit, Melanie Herschel, and Katerina Tzompanaki. 2014. Query-based why-not provenance with nedexplain. In *Extending database technology (EDBT)*.
- [12] Stefan Brass and Christian Goldberg. 2004. Detecting Logical Errors in SQL Queries. In *Tagungsband zum 16. GI-Workshop Grundlagen von Datenbanken, Mohnheim, NRW, Deutschland, 1.-4. Juni 2004*, Mireille Samia and Stefan Conrad (Eds.). Universität Düsseldorf, 28–32.
- [13] Stefan Brass and Christian Goldberg. 2005. Proving the Safety of SQL Queries. In *Fifth International Conference on Quality Software (QSIC 2005), 19-20 September 2005, Melbourne, Australia*. IEEE Computer Society, 197–204. <https://doi.org/10.1109/QSIC.2005.50>
- [14] Stefan Brass and Christian Goldberg. 2006. Semantic errors in SQL queries: A quite complete list. *J. Syst. Softw.* 79, 5 (2006), 630–644. <https://doi.org/10.1016/j.jss.2005.06.028>
- [15] Peter Buneman, Sanjeev Khanna, and Tan Wang-Chiew. 2001. Why and where: A characterization of data provenance. In *Database Theory—ICDT 2001: 8th International Conference London, UK, January 4–6, 2001 Proceedings* 8. Springer, 316–330.
- [16] Rafael Caballero, Yolanda García-Ruiz, and Fernando Sáenz-Pérez. 2012. Algorithmic debugging of SQL views. In *Perspectives of Systems Informatics: 8th International Andrei Ershov Memorial Conference, PSI 2011, Novosibirsk, Russia, June 27-July 1, 2011, Revised Selected Papers* 8. Springer, 77–85.
- [17] Rafael Caballero, Yolanda García-Ruiz, and Fernando Sáenz-Pérez. 2012. Declarative debugging of wrong and missing answers for SQL views. In *Functional and Logic Programming: 11th International Symposium, FLOPS 2012, Kobe, Japan, May 23-25, 2012. Proceedings* 11. Springer, 73–87.
- [18] Nofar Carmeli, Nikolaos Tziavelis, Wolfgang Gatterbauer, Benny Kimelfeld, and Mirek Riedewald. 2023. Tractable Orders for Direct Access to Ranked Answers of Conjunctive Queries. *ACM Trans. Database Syst.* 48, 1 (2023), 1:1–1:45. <https://doi.org/10.1145/3578517>
- [19] Nofar Carmeli, Shai Zeevi, Christoph Berkholz, Benny Kimelfeld, and Nicole Schweikardt. 2020. Answering (Unions of) Conjunctive Queries using Random Access and Random-Order Enumeration. In *Proceedings of the 39th ACM SIGMOD-SIGACT-SIGAI Symposium on Principles of Database Systems, PODS 2020, Portland, OR, USA, June 14–19, 2020*, Dan Suciu, Yufei Tao, and Zhewei Wei (Eds.). ACM, 393–409. <https://doi.org/10.1145/3375395.3387662>
- [20] Claudio Cerullo and Marco Porta. 2007. A system for database visual querying and query visualization: Complementing text and graphics to increase expressiveness. In *18th International Workshop on Database and Expert Systems Applications (DEXA 2007)*. IEEE, 109–113.
- [21] Bikash Chandra, Ananyo Banerjee, Udbhas Hazra, Mathew Joseph, and S. Sudarshan. 2021. Edit Based Grading of SQL Queries. In *CODS-COMAD 2021: 8th ACM IKDD CODS and 26th COMAD, Virtual Event, Bangalore, India, January 2-4, 2021*, Jayant R. Haritsa, Shourya Roy, Manish Gupta, Sharad Mehrotra, Balaji Vasan Srinivasan, and Yogesh Simmhan (Eds.). ACM, 56–64. <https://doi.org/10.1145/3430984.3431012>
- [22] Bikash Chandra, Bhupesh Chawda, Biplab Kar, K. V. Maheshwara Reddy, Shetal Shah, and S. Sudarshan. 2015. Data generation for testing and grading SQL queries. *Vldb J.* 24, 6 (2015), 731–755. <https://doi.org/10.1007/S00778-015-0395-0>
- [23] Adriane Chapman and HV Jagadish. 2009. Why not?. In *Proceedings of the 2009 ACM SIGMOD International Conference on Management of data*. 523–534.
- [24] Shumo Chu, Brendan Murphy, Jared Roesch, Alvin Cheung, and Dan Suciu. 2018. Axiomatic Foundations and Algorithms for Deciding Semantic Equivalences of SQL Queries. *Proc. VLDB Endow.* 11, 11 (2018), 1482–1495. <https://doi.org/10.14778/3236187.3236200>
- [25] Shumo Chu, Chenglong Wang, Konstantin Weitz, and Alvin Cheung. 2017. Cosette: An Automated Prover for SQL. In *8th Biennial Conference on Innovative Data Systems Research, CIDR 2017, Chaminade, CA, USA, January 8-11, 2017, Online Proceedings*. www.cidrdb.org. <http://cidrdb.org/cidr2017/papers/p51-chu-cidr17.pdf>
- [26] Shumo Chu, Konstantin Weitz, Alvin Cheung, and Dan Suciu. 2017. HoTTSQL: proving query rewrites with univalent SQL semantics. In *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2017, Barcelona, Spain, June 18-23, 2017*, Albert Cohen and Martin T. Vechev (Eds.). ACM, 510–524. <https://doi.org/10.1145/3062341.3062348>
- [27] Yingwei Cui, Jennifer Widom, and Janet L Wiener. 2000. Tracing the lineage of view data in a warehousing environment. *ACM Transactions on Database Systems (TODS)* 25, 2 (2000), 179–227.
- [28] Ralf Diestelkämper and Melanie Herschel. 2020. Tracing nested data with structural provenance for big data analytics. In *EDBT*. 253–264.
- [29] Benjamin Dietrich and Torsten Grust. 2015. A SQL debugger built from spare parts: Turning a SQL: 1999 database system into its own debugger. In *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data*. 865–870.
- [30] Kyriaki Dimitriadou, Olga Papaemmanouil, and Yanlei Diao. 2014. Explore-by-example: An automatic query steering framework for interactive data exploration. In *Proceedings of the 2014 ACM SIGMOD international conference on Management of data*. 517–528.
- [31] Haoran Ding, Zhaoguo Wang, Yicun Yang, Dexin Zhang, Zhenglin Xu, Haibo Chen, Ruzica Piskac, and Jinyang Li. 2023. Proving Query Equivalence Using Linear Integer Arithmetic. *Proc. ACM Manag. Data* 1, 4 (2023), 227:1–227:26. <https://doi.org/10.1145/3626768>
- [32] Idan Eldar, Nofar Carmeli, and Benny Kimelfeld. 2024. Direct Access for Answers to Conjunctive Queries with Aggregation. In *27th International Conference on Database Theory, ICDT 2024, March 25-28, 2024, Paestum, Italy (LIPIcs, Vol. 290)*, Graham Cormode and Michael Shekelyan (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 4:1–4:20. <https://doi.org/10.4230/LIPICS.ICDT.2024.4>
- [33] Full version of this submission. [n.d.]. https://anonymous.4open.science/r/oursys_eval-2063/. ([n.d.]).
- [34] Amir Gilad, Zhengjie Miao, Sudeepa Roy, and Jun Yang. 2022. Understanding Queries by Conditional Instances. In *SIGMOD '22: International Conference on Management of Data, Philadelphia, PA, USA, June 12 - 17, 2022*, Zachary G. Ives, Angela Bonifati, and Amr El Abbadi (Eds.). ACM, 355–368. <https://doi.org/10.1145/3514221.3517898>
- [35] Boris Glavic and Gustavo Alonso. 2009. Perm: Processing provenance and data on the same data model through query rewriting. In *2009 IEEE 25th International Conference on Data Engineering*. IEEE, 174–185.
- [36] Todd J Green, Grigoris Karvounarakis, and Val Tannen. 2007. Provenance semirings. In *Proceedings of the twenty-sixth ACM SIGMOD-SIGACT-SIGART symposium on Principles of database systems*. 31–40.
- [37] Torsten Grust and Jan Rittinger. 2013. Observing sql queries in their natural habitat. *ACM Transactions on Database Systems (TODS)* 38, 1 (2013), 1–33.
- [38] Laura M Haas, Johann Christoph Freytag, Guy M Lohman, and Hamid Pirahesh. 1989. Extensible query processing in Starburst. In *Proceedings of the 1989 ACM SIGMOD international conference on Management of data*. 377–388.
- [39] Sabaat Haroon, Chris Brown, and Muhammad Ali Gulzar. 2024. DeSQL: Interactive Debugging of SQL in Data-Intensive Scalable Computing. *Proc. ACM Softw. Eng.* 1, FSE (2024), 767–788. <https://doi.org/10.1145/3643761>
- [40] Yihao Hu, Amir Gilad, Kristin Stephens-Martinez, Sudeepa Roy, and Jun Yang. 2024. Qr-Hint: Actionable Hints Towards Correcting Wrong SQL Queries. *Proc. ACM Manag. Data* 2, 3 (2024), 164. <https://doi.org/10.1145/3654995>
- [41] Yihao Hu, Zhengjie Miao, Zhiming Leong, Haechan Lim, Zachary Zheng, Sudeepa Roy, Kristin Stephens-Martinez, and Jun Yang. 2022. I-Rex: An Interactive Relational Query Debugger for SQL. In *Proceedings of the 53rd ACM Technical Symposium on Computer Science Education V. 2*. 1180–1180.
- [42] Jiansheng Huang, Ting Chen, AnHai Doan, and Jeffrey F Naughton. 2008. On the provenance of non-answers to queries over extracted data. *Proceedings of the VLDB Endowment* 1, 1 (2008), 736–747.
- [43] Matteo Interlandi, Kshitij Shah, Sai Deep Tetali, Muhammad Ali Gulzar, Seunghyun Yoo, Miryung Kim, Todd Millstein, and Tyson Condie. 2015. Titian: Data provenance support in spark. In *Proceedings of the VLDB Endowment International Conference on Very Large Data Bases*, Vol. 9. NIH Public Access, 216.
- [44] Hannu Jaakkola and Bernhard Thalheim. 2003. Visual SQL-high-quality ER-based query treatment. In *Conceptual Modeling for Novel Application Domains: ER 2003 Workshops ECOMO, IWCQM, AOIS, and XSDM, Chicago, IL, USA, October 13, 2003. Proceedings* 22. Springer, 129–139.
- [45] Grigoris Karvounarakis, Todd J Green, Zachary G Ives, and Val Tannen. 2013. Collaborative data sharing via update exchange and provenance. *ACM Transactions on Database Systems (TODS)* 38, 3 (2013), 1–42.
- [46] Marie Le Guilly, Jean-Marc Petit, Vasilie-Marian Scuturici, and Ihab F Ilyas. 2019. Explique: Interactive databases exploration with SQL. In *Proceedings of the 28th ACM International Conference on Information and Knowledge Management*. 2877–2880.

- [47] Seokki Lee, Bertram Ludäscher, and Boris Glavic. 2019. PUG: a framework and practical implementation for why and why-not provenance. *The VLDB Journal* 28, 1 (2019), 47–71.
- [48] Aristotelis Leventidis, Jiahui Zhang, Cody Dunne, Wolfgang Gatterbauer, HV Jagadish, and Mirek Riedewald. 2020. QueryVis: Logic-based diagrams help users understand complicated SQL queries faster. In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*. 2303–2318.
- [49] Zhengjie Miao, Tiangang Chen, Alexander Bendeck, Kevin Day, Sudeepa Roy, and Jun Yang. 2020. I-Rex: an interactive relational query explainer for SQL. *Proceedings of the VLDB Endowment* 13, 12 (2020), 2997–3000.
- [50] Zhengjie Miao, Sudeepa Roy, and Jun Yang. 2019. Explaining Wrong Queries Using Small Examples. In *Proceedings of the 2019 International Conference on Management of Data, SIGMOD Conference 2019, Amsterdam, The Netherlands, June 30 - July 5, 2019*, Peter A. Boncz, Stefan Manegold, Anastasia Ailamaki, Amol Deshpande, and Tim Kraska (Eds.). ACM, 503–520. <https://doi.org/10.1145/3299869.3319866>
- [51] Daphne Miedema and George Fletcher. 2021. SQLVis: Visual query representations for supporting SQL learners. In *2021 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC)*. IEEE, 1–9.
- [52] Xing Niu, Boris Glavic, Ziyu Liu, Pengyuan Li, Dieter Gawlick, Vasudha Krishnaswamy, Zhen Hua Liu, and Danica Porobic. 2021. Provenance-based Data Skipping. *Proc. VLDB Endow.* 15, 3 (2021), 451–464. <https://doi.org/10.14778/3494124.3494130>
- [53] Laurel J. Orr, Srikanth Kandula, and Surajit Chaudhuri. 2019. Pushing Data-Induced Predicates Through Joins in Big-Data Clusters. *Proc. VLDB Endow.* 13, 3 (2019), 252–265. <https://doi.org/10.14778/3368289.3368292>
- [54] PostgreSQL. [n.d.]. <https://www.postgresql.org/>.
- [55] Kai Presler-Marshall, Sarah Heckman, and Kathryn T. Stolee. 2021. SQLRepair: Identifying and Repairing Mistakes in Student-Authored SQL Queries. In *43rd IEEE/ACM International Conference on Software Engineering: Software Engineering Education and Training, ICSE (SEET) 2021, Madrid, Spain, May 25-28, 2021*. IEEE, 199–210. <https://doi.org/10.1109/ICSE-SEET52601.2021.00030>
- [56] Fotis Psallidas and Eugene Wu. 2018. Smoke: Fine-grained lineage at interactive speed. *arXiv preprint arXiv:1801.07237* (2018).
- [57] Sivaprasad Sudhir, Wenbo Tao, Nikolay Pavlovich Laptev, Cyrille Habis, Michael J. Cafarella, and Samuel Madden. 2023. Pando: Enhanced Data Skipping with Logical Data Partitioning. *Proc. VLDB Endow.* 16, 9 (2023), 2316–2329. <https://doi.org/10.14778/3598581.3598601>
- [58] Nikolaos Tziavelis, Wolfgang Gatterbauer, and Mirek Riedewald. 2021. Beyond Equi-joins: Ranking, Enumeration and Factorization. *Proc. VLDB Endow.* 14, 11 (2021), 2599–2612. <https://doi.org/10.14778/3476249.3476306>
- [59] Margus Veanes, Nikolai Tillmann, and Jonathan de Halleux. 2010. Qex: Symbolic SQL Query Explorer. In *Logic for Programming, Artificial Intelligence, and Reasoning - 16th International Conference, LPAR-16, Dakar, Senegal, April 25-May 1, 2010, Revised Selected Papers (Lecture Notes in Computer Science, Vol. 6355)*, Edmund M. Clarke and Andrei Voronkov (Eds.). Springer, 425–446. https://doi.org/10.1007/978-3-642-17511-4_24
- [60] Shuxian Wang, Sicheng Pan, and Alvin Cheung. 2024. QED: A Powerful Query Equivalence Decider for SQL. *Proc. VLDB Endow.* 17, 11 (2024), 3602–3614. <https://www.vldb.org/pvldb/vol17/p3602-wang.pdf>
- [61] Cong Yan, Yin Lin, and Yeye He. 2023. Predicate Pushdown for Data Science Pipelines. *Proc. ACM Manag. Data* 1, 2 (2023), 136:1–136:28. <https://doi.org/10.1145/3589281>

A EXPERIMENT SUPPLEMENT

We present the remaining experiment results over TPC-H benchmark.

For each table in the TPC-H schema, we have the following indexes (all indexes are btree indexes in PostgreSQL):

- customer
 - primary index: c_custkey
 - secondary indexes: None
- lineitem
 - primary index: (l_orderkey, l_linenum)
 - secondary indexes: l_partkey, l_suppkey, l_shipdate
- nation
 - primary index: n_nationkey
 - secondary indexes: n_name, n_regionkey
- orders
 - primary index: o_orderkey
 - secondary indexes: o_custkey, o_orderdate
- part
 - primary index: p_partkey
 - secondary indexes: None
- partsupp
 - primary index: (ps_partkey, ps_suppkey)
 - secondary indexes: ps_suppkey
- region
 - primary index: r_regionkey
 - secondary indexes: None
- supplier
 - primary index: s_suppkey
 - secondary indexes: s_name, s_phone

We ran experiments for three different page size: 50, 100 and 200 for all tables (stages) in all TPC-H queries. For each table, we prepared three queries: milestone query, page query and table query and thus collecting the following data over all testing instances (i.e., 1GB, 5GB and 10GB):

- The execution time and output size of the milestone query and the table query.
- The execution time of the page query and baseline query (by rewriting the table query with OFFSET and LIMIT) for retrieving the first page (“head”), middle page (“mid”) and the second last (“tail”) page.

Since each query potentially contains multiple query blocks and subsequently multiple tables, we present the statistics for the largest table (measured in MB) to compute for each query. For page size 50, 100 and 200, the experiment results are shown in Table 4, Table 5 and Table 6 respectively.

In summary, we make the following conclusions:

- The milestone queries usually run slower than the table queries, but with acceptable delays since they are run at the beginning of the debugging session without affecting debugging operations later. In some cases, the milestone queries run faster than the table queries. On the other hand, the output sizes of the milestone queries are always smaller than those of the table queries by a rough factor of the page size.
- The optimization for page query almost always outperforms the baseline, and differences between the execution time grows as the database size grows, especially for “mid” and “tail” pages. The optimizations are sensitive to page size but insensitive to the database size. There are only few cases where the optimizations “over-hint” the PostgreSQL optimizer and cause the execution time to be roughly the same as the baseline.

Query	Page	1GB		5GB		10GB	
		Opt.	Base	Opt.	Base	Opt.	Base
Q1	head	0.151	0.053	0.065	0.057	0.067	0.148
	mid	0.156	448.025	0.123	14737.2125	0.122	31690.318
	tail	0.348	895.997	0.049	29474.368	0.143	63380.488
Q2	head	38.815	301.529	165.755	2167.54	368.719	4958.927
	mid	32.475	326.938	163.914	2168.14	364.631	5274.643
	tail	32.451	352.347	165.385	2168.74	359.418	5590.359
Q3	head	44.245	1380.082	12.04	82.834	12.396	88.617
	mid	5.416	1373.96	14.0	6705.6175	16.135	66430.58
	tail	4.618	1380.082	13.288	13328.401	18.388	132772.543
Q4	head	0.396	1.19	0.394	0.694	2.538	8.599
	mid	0.42	123.217	0.332	662.571	2.672	10941.241
	tail	0.44	245.244	0.933	1324.448	0.91	21873.883
Q5	head	9.79	14.0	14.181	16.402	221.318	115.96
	mid	9.987	1576.575	8.344	14288.552	91.998	41737.504
	tail	6.229	3139.15	7.667	28433.177	96.252	83359.048
Q6	head	0.605	0.244	0.799	0.689	0.61	0.73
	mid	0.727	277.99	1.061	1647.1675	2.566	44935.2465
	tail	0.522	555.743	1.564	3293.646	1.901	89556.745
Q7	head	28.103	17.004	31.474	24.972	590.231	69.64
	mid	17.029	346.029	30.229	4316.156	472.953	21205.945
	tail	641.349	675.054	20.782	8547.816	562.681	42342.25
Q8	head	13.237	7.953	13.711	17.009	147.984	26.324
	mid	11.648	1863.82	15.386	8362.616	117.491	25680.9055
	tail	13.777	3719.684	14.137	16708.223	141.207	51335.487
Q9	head	1.198	373.542	13.716	6431.233	15.416	29211.592
	mid	1.629	1610.2355	18.351	14402.325	13.01	46355.0735
	tail	1.149	2846.929	13.997	22099.616	12.7	63498.555
Q10	head	2.974	0.883	7.429	4.691	8.066	2.768
	mid	2.361	1036.81	5.482	11589.18	6.318	15474.274
	tail	1.751	2057.035	6.018	23173.671	4.664	30945.78
Q11	head	1.475	3.388	5.891	183.78	15.699	349.106
	mid	1.451	42.9335	5.189	463.9215	13.6	728.6065
	tail	1.331	82.479	5.3	744.063	14.108	1108.107
Q12	head	3.034	1.019	4.641	1.022	5.159	2.764
	mid	3.001	487.1765	2.9	8813.669	2.752	9562.74
	tail	1.017	973.334	3.207	17626.316	0.893	19122.716
Q13	head	0.26	0.189	0.224	0.199	0.227	0.57
	mid	0.241	473.5225	0.223	2631.6985	0.223	5854.7705
	tail	0.198	946.856	0.209	5263.198	0.22	11708.971
Q14	head	6.099	1.768	45.65	3.231	7.157	5.032
	mid	5.419	138.36	40.457	853.4155	4.334	1937.1275
	tail	5.484	274.959	39.291	1703.6	3.44	3869.223
Q15	head	0.402	0.282	1.17	0.369	1.195	0.281
	mid	0.47	203.9895	1.055	1309.741	3.017	2532.388
	tail	0.293	407.697	0.819	2619.113	2.122	5064.495
Q16	head	2.572	6.102	10.475	10.325	21.549	23.097
	mid	2.26	209.273	10.573	1253.456	21.849	2479.512
	tail	2.258	412.444	9.545	2496.587	20.773	4935.928
Q17	head	2.29	1.773	9.468	3.821	12.892	23.144
	mid	2.109	1349.49	4.934	11684.399	6.544	31238.4415
	tail	2.157	2697.224	4.672	23271.19	5.924	62453.739
Q18	head	24.219	3876.847	255.363	34380.398	10617.465	158664.235
	mid	25.32	6863.597	326.197	72862.803	9497.6	181742.5665
	tail	19.422	9850.347	244.225	111345.208	9660.102	204820.898
Q19	head	133.04	49.448	777.75	115.488	1350.159	851.97
	mid	141.193	95.092	781.639	449.903	1361.722	1106.9015
	tail	135.755	139.609	772.977	779.28	1346.508	1357.244
Q20	head	0.147	0.133	0.38	6288.547	0.245	0.137
	mid	0.114	83.62	0.407	8090.2775	0.868	1173.44
	tail	0.122	167.109	0.433	9892.008	0.719	2346.746
Q21	head	118.344	100.289	103.708	21934.138	466.694	53848.927
	mid	108.458	391.715	104.203	21757.629	384.542	53310.848
	tail	107.06	683.141	118.942	21581.122	511.79	53305.778
Q22	head	0.685	65.603	0.643	322.194	2.176	668.052
	mid	0.904	99.639	0.791	510.632	1.401	1034.1465
	tail	0.68	133.675	0.713	693.535	1.43	1400.241

(a) Optimization vs. Baseline for Page Query Execution Time

Query	DB size	Milestone Query		Table Query	
		Time (ms)	Output (MB)	Time (ms)	Output (MB)
Q1	1	10761.014	31.15872	2276.273	761.6532
	5	59026.436	155.85	32990.523	3809.74
	10	133380.121	311.66	66384.692	7618.51
Q2	1	269.273	0.0064	518.313	0.128
	5	1370.958	0.030	3681.9	0.646
	10	3140.467	0.062	7652.386	1.31
Q3	1	1355.772	0.216	2330.975	9.35844
	5	7515.589	1.055	57408.643	45.73
	10	20754.628	2.14	142657.569	92.93
Q4	1	381.742	0.382	584.866	6.36672
	5	1992.905	1.899	2151.989	31.649
	10	20009.641	3.79	38856.037	63.16
Q5	1	2616.26	0.935	3738.075	30.97
	5	13486.687	4.67	29643.753	154.836
	10	36375.305	9.324	90812.433	308.884
Q6	1	982.034	1.79	1002.881	43.884
	5	4977.102	8.95	3831.173	218.803
	10	16230.004	17.923	93536.412	438.125
Q7	1	717.982	0.151	1546.118	0.947
	5	8766.523	0.73	9575.02	4.602
	10	18707.733	1.48	43991.291	9.306
Q8	1	5219.431	11.66	4740.349	65.639
	5	30334.18	58.05	21427.097	326.54
	10	64616.043	116.28	100710.36	654.091
Q9	1	2583.215	4.26	3782.709	23.998
	5	20082.296	20.9	23844.223	118.06
	10	41861.057	42.011	65514.842	236.31
Q10	1	2116.497	2.89	2676.497	77.86
	5	10989.395	14.53	12381.126	390.69
	10	24260.84	29.047	38998.348	780.63
Q11	1	193.031	0.21	212.718	2.10
	5	630.802	1.15	730.264	28.145
	10	1333.583	2.298	15813.967	56.193
Q12	1	1629.965	0.788	2175.2	10.506
	5	6427.45	3.94	22007.229	52.62
	10	14693.272	7.879	27017.67	105.057
Q13	1	5126.519	7.192	2107.894	74.4
	5	25697.736	35.96	11627.507	372.00
	10	58677.526	71.92	25327.205	744.0
Q14	1	346.662	0.539	456.689	18.57
	5	1925.3	92.79	7269.705	2.69
	10	3981.041	5.391	35628.615	185.704
Q15	1	888.715	2.16	682.545	52.91
	5	4503.56	10.84	9401.36	265.098
	10	10777.325	21.696	39915.163	530.346
Q16	1	606.336	0.74	508.36	16.45
	5	4606.968	3.678	4861.268	81.739
	10	5761.204	7.377	6919.425	163.94
Q17	1	3848.482	1.44	3438.054	11.2
	5	21498.169	7.2	25793.157	56.0
	10	57233.481	14.4	65891.692	112.0
Q18	1	53581.528	52.21	17478.406	1272.73
	5	267399.772	260.86	136289.565	6358.46
	10	255628.375	521.39	257795.588	12709.09
Q19	1	157.224	0.043	401.107	1.50
	5	930.834	0.217	8263.581	7.47
	10	1559.524	0.421	35581.941	14.506
Q20	1	2779.138	6.57	2549.86	1056.21
	5	13674.114	32.885	49759.313	5279.963
	10	35559.011	65.69	97257.186	10557.545
Q21	1	816.355	0.0057	1422.404	0.0227
	5	13747.978	0.249	26823.763	6.538
	10	35719.077	0.511	94891.68	13.41
Q22	1	156.275	0.045	159.611	0.404
	5	773.757	0.228	869.385	2.035
	10	1573.278	0.457	1975.12	4.069

(b) Milstone Query vs. Table Query

Table 4: Experiment results for Page Size 50

Query	Page	1GB		5GB		10GB	
		Opt.	Base	Opt.	Base	Opt.	Base
Q1	head	0.091	0.081	0.301	0.234	0.113	0.092
	mid	0.086	426.515	0.441	15051.03	0.812	30014.521
	tail	0.076	852.949	0.222	30101.825	0.078	60028.951
Q2	head	306.888	308.865	2022.056	2393.282	4850.654	4851.175
	mid	292.801	333.994	2008.593	2521.091	4877.578	5213.552
	tail	293.259	359.122	1984.249	2648.899	4985.901	5575.929
Q3	head	206.687	6.743	20.239	79.53	20.263	140.854
	mid	1.519	333.901	18.553	7828.453	12.9	22104.909
	tail	1.61	661.058	17.36	15494.749	14.079	43262.151
Q4	head	0.645	0.832	0.697	0.889	0.709	2.187
	mid	0.774	119.444	1.179	702.712	1.392	12318.296
	tail	0.72	238.055	1.204	1404.536	1.263	24415.675
Q5	head	18.027	15.82	24.736	20.695	181.443	54.614
	mid	17.306	1562.262	18.946	18932.159	120.167	44299.05
	tail	13.122	3108.703	18.767	37843.623	117.645	88045.168
Q6	head	1.116	0.489	1.903	1.699	4.542	0.538
	mid	1.037	279.391	1.718	1591.986	4.372	42536.88
	tail	0.82	558.293	2.68	3182.272	3.147	84384.737
Q7	head	629.443	39.648	64.844	51.109	924.415	56.157
	mid	625.804	352.967	51.26	4280.98	601.548	19926.793
	tail	634.1	664.976	35.191	8510.85	573.285	39797.43
Q8	head	27.424	11.62	28.141	22.043	238.584	39.353
	mid	22.1	1853.264	30.463	8381.889	230.936	37035.365
	tail	23.12	3694.909	28.479	16741.734	248.284	74031.377
Q9	head	2.227	377.075	25.992	6366.156	27.88	28444.105
	mid	1.941	1615.231	34.151	14037.594	25.205	46211.428
	tail	2.05	2853.388	23.469	21709.033	25.557	63978.751
Q10	head	1.974	2.556	9.511	12.751	10.925	3.398
	mid	1.619	1026.081	9.452	11387.6	10.205	15272.071
	tail	1.854	2049.606	9.508	22762.449	7.526	30540.744
Q11	head	54.254	4.695	299.224	10.559	619.342	11.336
	mid	54.957	45.489	318.275	249.877	724.327	494.175
	tail	39.354	80.179	315.443	489.194	558.471	977.013
Q12	head	5.631	4.889	3.996	1.703	8.291	1.944
	mid	6.348	498.536	2.223	8768.324	2.118	9526.046
	tail	1.837	992.183	1.927	17534.945	1.554	19050.148
Q13	head	0.386	0.248	0.781	0.283	0.391	0.352
	mid	0.394	469.518	0.777	2591.396	0.348	5733.727
	tail	0.344	938.788	0.768	5182.509	0.348	11467.101
Q14	head	6.608	3.691	41.436	5.013	130.602	9.195
	mid	5.8	140.401	41.312	936.271	124.42	2149.62
	tail	6.099	277.112	40.986	1867.53	119.398	4290.046
Q15	head	0.68	0.475	1.987	1.38	3.843	0.503
	mid	0.67	205.531	3.212	1217.152	3.615	2482.398
	tail	0.773	410.586	1.202	2432.924	2.257	4964.293
Q16	head	2.572	2.434	13.604	14.062	21.648	23.848
	mid	2.971	205.463	11.034	1186.252	22.334	2424.724
	tail	2.699	408.492	10.095	2358.442	21.343	4825.6
Q17	head	4.207	3.125	9.093	4.079	19.093	3.759
	mid	4.017	1333.833	10.083	11583.229	14.075	29799.824
	tail	4.054	2664.541	9.281	23162.379	28.103	59595.89
Q18	head	24.944	4339.022	5120.907	75225.846	46208.018	170402.011
	mid	25.239	7085.561	5572.922	86029.494	8721.564	193398.846
	tail	20.279	9832.101	5314.759	96833.142	10349.474	216395.681
Q19	head	131.08	95.883	785.001	187.103	1355.466	851.254
	mid	130.406	120.139	766.274	490.312	1353.755	1100.595
	tail	128.953	136.898	773.398	793.522	1378.008	1349.937
Q20	head	0.252	0.202	0.3	0.197	0.404	1.248
	mid	0.25	86.114	1.002	466.788	1.744	1149.965
	tail	0.216	172.025	0.927	933.379	1.024	2298.681
Q21	head	934.343	802.551	12261.892	18646.15	391.055	52733.088
	mid	809.447	802.856	4749.522	18599.982	386.668	52339.251
	tail	800.625	803.161	4779.119	18646.15	330.754	52221.527
Q22	head	1.18	64.627	4.756	350.916	1.25	671.538
	mid	1.369	99.096	3.578	525.755	1.324	1034.214
	tail	1.212	133.565	3.283	700.593	3.717	1395.944

(a) Optimization vs. Baseline for Page Query Execution Time

Query	DB size	Milestone Query		Table Query	
		Time (ms)	Output (MB)	Time (ms)	Output (MB)
Q1	1	10784.539	15.579	2491.228	761.653
	5	65880.08	77.927	32098.349	3809.747
	10	130708.242	155.834	63821.418	7618.516
Q2	1	269.154	0.003	499.973	0.129
	5	1552.798	0.015	3618.162	0.647
	10	3199.822	0.031	7966.225	1.31
Q3	1	1380.654	0.108	2328.746	9.358
	5	7473.251	0.528	61037.274	45.735
	10	16236.185	1.072	45154.879	92.932
Q4	1	368.092	0.191	578.802	6.367
	5	1943.838	0.95	2155.429	31.649
	10	19967.007	1.895	32459.15	63.165
Q5	1	2639.434	0.468	3757.791	30.979
	5	17407.391	2.337	39228.976	154.836
	10	39188.193	4.662	91746.208	308.884
Q6	1	996.277	0.898	981.971	43.884
	5	5624.262	4.476	3881.74	218.804
	10	16990.111	8.962	70120.717	438.125
Q7	1	709.801	0.076	1390.292	0.947
	5	8878.686	0.369	9633.558	4.602
	10	19075.358	0.745	41615.446	9.306
Q8	1	5231.564	5.835	4653.983	65.639
	5	30111.784	29.027	21438.65	326.55
	10	81725.932	58.142	116668.268	654.091
Q9	1	2570.109	2.134	3691.996	23.998
	5	22531.648	10.495	24503.937	118.061
	10	41083.291	21.006	63760.976	236.312
Q10	1	2130.495	1.449	2626.755	77.861
	5	11754.423	7.269	12304.497	390.698
	10	23769.229	14.524	38235.373	780.631
Q11	1	131.609	0.121	128.167	5.928
	5	702.067	0.576	890.554	28.146
	10	1366.91	1.149	15656.266	56.193
Q12	1	1637.115	0.394	2208.423	10.507
	5	6712.775	1.974	21493.059	52.627
	10	14757.48	3.94	26192.762	105.057
Q13	1	5091.584	3.596	2080.932	74.4
	5	28848.33	17.98	11843.681	372.001
	10	59175.603	35.96	24853.255	744.001
Q14	1	336.445	0.27	448.451	18.571
	5	2061.857	1.347	7470.676	92.793
	10	4240.96	2.696	35069.773	185.705
Q15	1	919.933	1.083	680.466	52.914
	5	5394.991	5.423	10751.989	265.099
	10	10649.701	10.848	38686.0	530.346
Q16	1	537.092	0.37	510.334	16.455
	5	2824.758	1.839	4811.792	81.74
	10	5547.766	3.689	6748.268	163.946
Q17	1	3773.283	0.72	3388.614	11.2
	5	27777.519	3.6	30992.322	56.0
	10	58270.198	7.2	66574.781	112.0
Q18	1	39429.388	26.108	17182.441	1272.737
	5	125175.085	130.43	127113.513	6358.46
	10	253399.157	260.7	273304.685	12709.096
Q19	1	159.099	0.022	391.195	1.501
	5	953.632	0.109	5059.877	7.47
	10	1567.009	0.211	37647.341	14.506
Q20	1	2764.728	3.289	2614.205	1056.214
	5	15222.367	16.443	42226.478	5279.964
	10	36583.623	32.845	100012.771	10557.545
Q21	1	817.198	0.026	1410.427	1.332
	5	13901.619	0.125	29317.22	6.538
	10	35891.359	0.256	56547.062	13.417
Q22	1	152.344	0.023	159.713	0.405
	5	809.068	0.114	970.405	2.035
	10	1559.652	0.229	1702.751	4.069

(b) Milstone Query vs. Table Query

Table 5: Experiment results for Page Size 100

Query	Page	1GB		5GB		10GB	
		Opt.	Base	Opt.	Base	Opt.	Base
Q1	head	0.143	0.386	0.43	0.208	0.21	0.396
	mid	0.138	437.834	0.466	16745.148	0.215	30184.466
	tail	0.128	875.282	0.386	33490.088	0.151	60368.535
Q2	head	298.866	306.012	2097.534	2480.709	4404.559	5117.611
	mid	293.025	339.097	2119.439	2640.438	4623.84	5470.365
	tail	292.456	372.183	2114.01	2800.167	4415.682	5823.119
Q3	head	228.084	12.503	34.568	88.792	39.79	223.527
	mid	165.074	346.972	24.108	8011.728	46.772	59672.52
	tail	161.195	681.441	31.033	15934.664	44.362	119061.532
Q4	head	3.966	4.467	1.518	1.881	7.457	20.433
	mid	4.213	125.054	2.239	713.966	10.155	11032.907
	tail	2.707	245.641	2.823	1425.922	10.793	21861.043
Q5	head	29.002	29.987	34.971	73.538	526.098	433.531
	mid	26.733	1600.75	38.204	19050.377	298.907	42043.691
	tail	23.722	3171.513	37.134	37859.367	269.094	82350.115
Q6	head	1.351	1.048	4.533	1.106	4.22	1.223
	mid	1.362	276.844	1.415	1614.682	5.425	41153.139
	tail	1.29	552.641	2.156	3228.257	3.205	82305.056
Q7	head	646.875	67.955	121.792	84.835	1470.09	80.1
	mid	637.23	388.011	8776.295	4399.466	911.324	19518.046
	tail	635.496	675.246	8789.561	8714.097	878.456	38955.992
Q8	head	54.945	18.989	51.474	26.841	371.484	55.388
	mid	45.471	1890.527	56.306	8322.906	345.375	25782.285
	tail	44.974	3762.065	50.205	16618.97	341.093	51509.182
Q9	head	3.896	376.365	41.004	6194.939	48.271	2917.899
	mid	3.875	1636.117	47.682	13998.14	46.527	46923.13
	tail	3.295	2862.053	48.551	21801.34	48.611	64728.361
Q10	head	3.48	2.143	17.37	6.903	25.062	5.472
	mid	3.311	1013.409	12.885	11619.602	20.801	15413.302
	tail	3.665	2024.675	10.083	23232.301	20.876	30821.132
Q11	head	73.072	11.056	571.28	14.588	1178.938	18.543
	mid	33.226	45.024	575.075	253.133	1246.811	516.043
	tail	29.083	78.992	544.629	491.679	1106.774	1013.543
Q12	head	8.877	3.224	4.726	3.992	5.459	3.697
	mid	8.334	481.152	4.094	9102.369	4.591	9516.675
	tail	3.599	959.08	3.62	18200.746	3.684	19029.654
Q13	head	0.659	1.076	0.867	0.495	1.387	0.495
	mid	0.618	491.226	0.765	2654.764	1.358	5627.7
	tail	0.57	981.375	0.947	6597.823	0.562	11254.905
Q14	head	7.721	5.937	45.371	15.041	132.967	13.395
	mid	6.923	139.767	46.699	885.645	119.878	1904.49
	tail	6.915	273.596	45.787	1734.791	123.092	3795.585
Q15	head	1.299	0.871	3.158	2.645	8.039	0.934
	mid	1.056	204.815	10.181	1307.789	5.555	2541.68
	tail	1.222	408.76	3.237	2612.933	7.332	5082.425
Q16	head	6.448	2.53	11.776	11.327	22.839	30.073
	mid	3.783	209.971	12.759	1266.214	30.676	2455.84
	tail	10.755	414.208	11.599	2521.1	23.023	4881.606
Q17	head	17.165	6.519	22.959	7.215	34.088	8.65
	mid	8.918	1233.832	22.988	16872.764	40.368	28547.251
	tail	8.324	2461.146	22.979	33738.314	29.178	57085.851
Q18	head	2379.584	2916.801	21528.099	81351.562	48625.669	169993.452
	mid	1814.595	6061.662	19208.862	96954.308	48567.652	192353.968
	tail	2723.067	9206.524	23889.779	112557.054	51898.173	214714.485
Q19	head	135.689	168.863	768.251	620.145	1326.267	1174.615
	mid	134.435	154.711	765.594	694.52	1336.209	1263.101
	tail	130.649	142.62	778.368	768.895	1335.71	1339.853
Q20	head	0.469	0.146	1.302	0.12	2.258	0.742
	mid	0.429	84.42	1.228	557.974	2.838	1117.773
	tail	0.406	168.695	1.128	1115.828	2.171	2234.804
Q21	head	975.017	826.083	5911.734	20059.439	122272.423	62562.862
	mid	812.574	821.873	4663.958	19896.141	28701.653	57648.24
	tail	797.346	821.202	4649.823	19732.842	18877.568	52733.618
Q22	head	7.382	67.275	3.595	334.122	3.83	672.021
	mid	2.568	101.09	3.72	512.774	3.418	1036.466
	tail	2.384	133.957	3.3	691.427	2.992	1383.612

(a) Optimization vs. Baseline for Page Query Execution Time

Query	DB size	Milestone Query		Table Query	
		Time (ms)	Output (MB)	Time (ms)	Output (MB)
Q1	1	12489.981	7.79	2290.842	761.653
	5	65474.109	38.964	38578.118	3809.747
	10	132658.807	77.917	66244.523	7618.516
Q2	1	278.907	0.003	508.325	0.129
	5	1625.016	0.008	3571.426	0.647
	10	3228.015	0.016	7803.155	1.31
Q3	1	1361.843	0.054	2337.962	9.358
	5	7514.865	0.264	35611.169	45.735
	10	20991.536	0.536	120721.778	92.932
Q4	1	369.552	0.096	589.573	6.367
	5	1933.858	0.475	2977.123	31.649
	10	19781.825	0.948	37888.832	63.165
Q5	1	2679.392	0.234	3930.441	30.979
	5	17689.275	1.169	40644.4	154.836
	10	36275.997	2.332	90122.184	308.884
Q6	1	1038.021	0.449	1013.535	43.884
	5	5707.782	2.238	4731.004	218.804
	10	16225.613	4.481	83066.905	438.125
Q7	1	719.692	0.038	1561.921	0.947
	5	8946.554	0.184	10968.467	4.602
	10	18809.911	0.372	40687.927	9.306
Q8	1	5889.245	2.918	4722.149	65.639
	5	30405.964	14.514	22748.508	326.55
	10	61301.514	29.071	93827.029	654.091
Q9	1	2667.144	1.067	3857.448	23.998
	5	22619.108	5.247	25742.462	118.061
	10	41691.015	10.503	66079.948	236.312
Q10	1	2163.786	0.724	2663.698	77.861
	5	11796.093	3.635	13510.529	390.698
	10	24423.312	7.262	38940.872	780.631
Q11	1	134.791	0.061	127.976	5.928
	5	708.581	0.288	859.238	28.146
	10	1320.652	0.575	15533.24	56.193
Q12	1	1643.257	0.197	2236.642	10.507
	5	6881.72	0.987	21892.578	52.627
	10	14647.287	1.97	26779.409	105.057
Q13	1	5727.408	1.798	2146.298	74.4
	5	29358.318	8.59	12564.698	372.001
	10	58707.211	17.98	25154.396	744.001
Q14	1	366.078	0.135	456.921	18.571
	5	2023.895	0.674	18845.022	92.793
	10	4006.049	1.348	35463.586	185.705
Q15	1	967.879	0.541	683.278	52.914
	5	5233.873	2.712	20492.43	265.099
	10	10787.142	5.424	40126.643	530.346
Q16	1	516.711	0.185	503.538	16.455
	5	2663.285	0.92	3373.281	81.74
	10	5574.077	1.845	6639.659	163.946
Q17	1	5936.046	0.36	5918.6	11.2
	5	28270.684	1.8	34331.977	56.0
	10	57344.405	3.6	65566.539	112.0
Q18	1	20666.734	13.054	16486.628	1272.737
	5	117514.027	65.215	143138.813	6358.46
	10	246810.647	130.35	270944.734	12709.096
Q19	1	158.533	0.011	396.758	1.501
	5	937.639	0.054	10230.62	7.47
	10	1518.954	0.105	43848.411	14.506
Q20	1	3001.687	1.644	2566.601	1056.214
	5	15469.147	8.221	49269.443	5279.964
	10	35479.613	16.423	100506.581	10557.545
Q21	1	813.243	0.013	1457.584	1.332
	5	14187.38	0.063	25798.959	6.538
	10	35394.015	0.128	95884.127	13.417
Q22	1	157.132	0.012	159.901	0.405
	5	779.235	0.057	864.045	2.035
	10	1572.976	0.114	1940.377	4.069

(b) Milstone Query vs. Table Query

Table 6: Experiment results for Page Size 200

B USER STUDY SUPPLEMENT

The user study is conducted with an instance of beer database with the following schema (keys are underlined):

- Drinker (name, address)
- Bar (name, address)
- Beer (name, brewery)
- Frequents (drinker, bar, times_a_week)
- Serves (bar, beer, price)
- Likes (drinker, beer)

B.1 Debugging Quesiton 1

Question Statement. For each bar Ben visits, find price of the most expensive and cheapest drink at that bar. Format the output as (bar, price), no duplicates.

The wrong query presented to the students are as follows:

```
|| WITH t1 AS (  
||   SELECT bar, price  
||   FROM serves  
||   WHERE price = (  
||     SELECT MAX(S1.price)  
||     FROM serves S1  
||     WHERE S1.bar = bar  
||   )  
||   UNION ALL  
||   SELECT bar, price  
||   FROM serves  
||   WHERE price = (  
||     SELECT MIN(S1.price)  
||     FROM serves S1  
||     WHERE S1.bar = bar  
||   )  
|| )  
|| SELECT t1.bar, t1.price  
|| FROM t1, frequents  
|| WHERE t1.bar = frequents.bar
```

```
||   AND frequents.drinker = 'Ben';
```

The above query has two mistakes:

- UNION ALL creates duplicates when the most expensive and cheapest drink share the same price.
- The bar in both scalar subqueries are referencing the wrong column. Without correct aliasing, both bar refer to the bar in S1, making the WHERE condition a tautology.

B.2 Debugging Quesiton 2

Question Statement. Suppose every time a drinker frequents a bar, he buys all his favorite beers at that bar. Find the expected weekly revenue of each bar and rank them by the revenue from high to low. The output should be in the format of (bar, revenue). If a bar is not frequented by any drinker, or it does not serve any beer, or none of its beer is liked by any drinker, output (bar, NULL).

The wrong query presented to the students are as follows:

```
|| SELECT S.bar,  
||   SUM(F.times_a_week * SUM(S.price) AS revenue  
|| FROM serves S,  
||   frequents F,  
||   likes L  
|| WHERE S.bar = F.bar  
||   AND S.beer = L.beer  
|| GROUP BY S.bar  
|| ORDER BY revenue DESC;
```

The above query has three mistakes:

- The join predicate F.drinker = L.drinker is missing.
- The expression for sum is incorrect as it will blow up the result. The correct expression is SUM(F.times_a_week * S.price).
- There will be no “NULL” tuple produced by the query, i.e., bars which do not serve any beer / serve no liked by anyone will not be included in the result.